## 

```
(https://databricks.com)
   from pyspark.sql import SparkSession
   spark = SparkSession.builder.appName('lin_reg').getOrCreate()
   from pyspark.ml.regression import LinearRegression
   # Load Dataset
   df = spark.read.format("csv").option("header", "true").load("dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/Linear_regression_d
   # validate the size of the data
   print((df.count(), len(df.columns)))
 (1232, 6)
   # explore the data
   df.printSchema()
  |-- var_1: integer (nullable = true)
  |-- var_2: integer (nullable = true)
  |-- var_3: integer (nullable = true)
  |-- var_4: double (nullable = true)
  |-- var_5: double (nullable = true)
  |-- output: double (nullable = true)
   # view statistical mesures of the data
   df.describe().show(5,False)
 +-----+
 |count |1232 |1232 |1232
                                                         1232
                                                                            1232
                                                                                               1232
         \lceil 715.0819805194806 \rceil 715.0819805194806 \rceil 80.90422077922078 \ \lceil 0.3263311688311693 \ \rceil \ 0.2592727272727275 \ \lceil 0.39734172077922014 \rceil 
 |stddev |91.5342940441652 |93.07993263118064|11.458139049993724|0.015012772334166148|0.012907228928000298|0.03326689862173776|
              |472
                                |40
                                                0.277
 |min |463
                                                                    0.214
 |max | 1009
                        1103
                                        |116
                                                          0.373
                                                                            0.294
                                                                                                0.491
   # sneak into the dataset
   df.head(3)
 Out[78]: [Row(var_1=734, var_2=688, var_3=81, var_4=0.328, var_5=0.259, output=0.418),
  Row(var_1=700, var_2=600, var_3=94, var_4=0.32, var_5=0.247, output=0.389),
  Row(var_1=712, var_2=705, var_3=93, var_4=0.311, var_5=0.247, output=0.417)]
   # import corr function from pyspark functions
   from pyspark.sql.functions import corr
   # check for correlation & how to add variable using corr functions
   df.select(corr('var_1', 'output')).show()
```

```
|corr(var_1, output)|
| 0.9187399607627283|
+----+
  # import
  from pyspark.ml.linalg import Vector
  from pyspark.ml.feature import VectorAssembler
  # select the columns to create the input vector
  df.columns
Out[82]: ['var_1', 'var_2', 'var_3', 'var_4', 'var_5', 'output']
  # create the vector assembler
  vec_assembler = VectorAssembler(inputCols = ['var_1', 'var_2', 'var_3', 'var_4', 'var_5'], outputCol='features')
  # transform
  features_df = vec_assembler.transform(df)
  features_df.printSchema()
root
 |-- var_1: integer (nullable = true)
 |-- var_2: integer (nullable = true)
 |-- var_3: integer (nullable = true)
 |-- var_4: double (nullable = true)
 |-- var_5: double (nullable = true)
 |-- output: double (nullable = true)
 |-- features: vector (nullable = true)
  features_df.select('features').show(5,False)
features
|[734.0,688.0,81.0,0.328,0.259]|
|[700.0,600.0,94.0,0.32,0.247] |
|[712.0,705.0,93.0,0.311,0.247]|
|[734.0,806.0,69.0,0.315,0.26] |
|[613.0,759.0,61.0,0.302,0.24] |
+----+
only showing top 5 rows
  model_df = features_df.select('features', 'output')
  model_df.show(5,False)
```

```
|features |output|
|[734.0,688.0,81.0,0.328,0.259]|0.418 |
|[700.0,600.0,94.0,0.32,0.247] |0.389 |
|[712.0,705.0,93.0,0.311,0.247]|0.417 |
|[734.0,806.0,69.0,0.315,0.26] |0.415 |
|[613.0,759.0,61.0,0.302,0.24] |0.378 |
only showing top 5 rows
   print((model_df.count(), len(model_df.columns)))
 (1232, 2)
Spilt data - Train & Test Data
   # sec
   train_df, test_df = model_df.randomSplit([0.7,0.3])
   print((train_df.count(), len(train_df.columns)))
(880, 2)
   print((test_df.count(), len(test_df.columns)))
(352, 2)
```

train\_df.describe().show()

# Build Linear Regression Model

lin\_reg = LinearRegression(labelCol='output')

```
# fit the linear model on training dataset
  lr_model = lin_reg.fit(train_df)
  lr_model.intercept
Out[100]: 0.18626961441250367
  print(lr_model.coefficients)
[0.0003371220458572444, 5.991498706752434 e-05, 0.0002556814872303228, -0.6686443805619972, 0.48044283582971575]
  training_predictions = lr_model.evaluate(train_df)
  {\tt training\_predictions.meanSquaredError}
Out[103]: 0.00013704630021768094
  training_predictions.r2
Out[104]: 0.8772919740056858
  # make predictions on the test data
  test_results= lr_model.evaluate(test_df)
  # view the residual errors based on the predictions
  test_results.residuals.show(10)
+----+
         residuals
[0.009429260836439135]
|-9.42564848742444...|
|0.013713981224977412|
|-0.01210437809228...|
|-0.00679164860545...|
|0.010048355854590463|
| 8.97412483862492E-4|
I -0.012396549858743|
|-0.01181377619641...|
|-0.00117635064512...|
only showing top 10 rows
  # coefficient of determination value for model
  test_results.r2
Out[107]: 0.8476131884871057
```

 ${\tt test\_results.rootMeanSquaredError}$ 

Out[108]: 0.012796007916146983

Out[109]: 0.00016373781859009624