

```
(https://databricks.com)
   # create spark session
   from pyspark.sql import SparkSession
   spark = SparkSession.builder.appName('nlp').getOrCreate()
   df = spark.createDataFrame([(1, 'I Really Liked this movie'),
                           (2, 'I Would recommend this movie to my friends'),
                           (3, 'movie was alright but acting was horrible'),
                           (4, 'I am never watching that movie ever again')],
                          ['user_id', 'review'])
   df.show(5,False)
 +-----+
 |user_id|review
       |I Really Liked this movie
 12
       |I Would recommend this movie to my friends|
 |3
      |movie was alright but acting was horrible |
 |4 |I am never watching that movie ever again |
   # tokenization
   from pyspark.ml.feature import Tokenizer
   tokenization = Tokenizer(inputCol='review', outputCol='tokens')
   tokenized_df = tokenization.transform(df)
   tokenized_df.show(4,False)
 |user_id|review
                                             tokens
 1
       [i, really, liked, this, movie]
 |2
        |I Would recommend this movie to my friends|[i, would, recommend, this, movie, to, my, friends]|
 13
        |movie was alright but acting was horrible |[movie, was, alright, but, acting, was, horrible] |
       |I am never watching that movie ever again |[i, am, never, watching, that, movie, ever, again] |
 +-----
   # Stopwords removal
   from pyspark.ml.feature import StopWordsRemover
```

```
stopword_removal = StopWordsRemover(inputCol='tokens', outputCol='refined_tokens')
 refined_df = stopword_removal.transform(tokenized_df)
 refined_df.select(['user_id', 'tokens', 'refined_tokens']).show(10,False)
|user_id|tokens
                                                refined_tokens
+-----+
1
     [i, really, liked, this, movie]
                                               [really, liked, movie]
    [i, would, recommend, this, movie, to, my, friends]|[recommend, movie, friends]
13
    |[movie, was, alright, but, acting, was, horrible] |[movie, alright, acting, horrible]|
     |[i, am, never, watching, that, movie, ever, again] |[never, watching, movie, ever] |
4
+-----+
 # Count Vectorizer
 from pyspark.ml.feature import CountVectorizer
 count_vec = CountVectorizer(inputCol='refined_tokens', outputCol='features')
 cv_df= count_vec.fit(refined_df).transform(refined_df)
 cv_df.select(['user_id', 'refined_tokens', 'features']).show(4, False)
+-----+
|user_id|refined_tokens
                                 features
+-----+-----
|1 | [really, liked, movie] | (11,[0,3,7],[1.0,1.0,1.0])
|2 | [recommend, movie, friends] | (11,[0,6,9],[1.0,1.0,1.0])
     |[movie, alright, acting, horrible]|(11,[0,2,4,5],[1.0,1.0,1.0,1.0]) |
|3
     |[never, watching, movie, ever] |(11,[0,1,8,10],[1.0,1.0,1.0])|
4
 count_vec.fit(refined_df).vocabulary
Out[23]: ['movie',
 'liked',
 'horrible',
 'friends',
 'ever'.
 'alright',
 'recommend',
 'acting',
 'really',
 'never',
 'watching']
```

```
# tf_idf
 from pyspark.ml.feature import HashingTF,IDF
 hashing_vec = HashingTF(inputCol="refined_tokens", outputCol='tf_features')
 hashing_df = hashing_vec.transform(refined_df)
 hashing_df.select(['user_id', 'refined_tokens', 'tf_features']).show(4, False)
|user_id|refined_tokens
+-----+
   |[really, liked, movie] |(262144,[99172,210223,229264],[1.0,1.0,1.0])
|[recommend, movie, friends] |(262144,[68228,130047,210223],[1.0,1.0,1.0])
1
|2
      |[movie, alright, acting, horrible]|(262144,[95685,171118,210223,236263],[1.0,1.0,1.0,1.0])|
|3
4
      |[never, watching, movie, ever] |(262144,[63139,113673,203802,210223],[1.0,1.0,1.0,1.0])|
+-----+
 tf_idf_vec = IDF(inputCol='tf_features', outputCol='tf_idf_features')
 tf_idf_df = tf_idf_vec.fit(hashing_df).transform(hashing_df)
 tf_idf_df.select(['user_id', 'tf_idf_features']).show(4, False)
+-----+
| {\tt user\_id}| {\tt tf\_idf\_features}
11
      |(262144,[99172,210223,229264],[0.9162907318741551,0.0,0.9162907318741551])
      [(262144,[68228,130047,210223],[0.9162907318741551,0.9162907318741551,0.0])
|3
      |(262144,[95685,171118,210223,236263],[0.9162907318741551,0.9162907318741551,0.9162907318741551])|
      [(262144,[63139,113673,203802,210223],[0.9162907318741551,0.9162907318741551,0.9162907318741551,0.0])
 # classification
 text_df = spark.read.csv('dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/Movie_reviews.csv', inferSchema=True,
 header=True, sep=',')
 text df.printSchema()
```

```
|-- Review: string (nullable = true)
|-- Sentiment: string (nullable = true)
  text_df.count()
Out[36]: 7087
  from pyspark.sql.functions import rand
  text_df.orderBy(rand()).show(10,False)
Review
                                                                 |Sentiment|
|Harry Potter is AWESOME I don't care if anyone says differently!..
                                                                 1
|Harry Potter is AWESOME I don't care if anyone says differently!..
                                                                 |1
|Always knows what I want, not guy crazy, hates Harry Potter..
                                                                 10
|I either LOVE Brokeback Mountain or think it's great that homosexuality|1
|dudeee i LOVED brokeback mountain!!!!
                                                                 |1
II liked the movie Brokeback Mountain.
                                                                 11
|I love Harry Potter.
                                                                  1
|i hate brokeback mountain!!
                                                                 10
Brokeback Mountain was an AWESOME movie.
                                                                 1
\mid Harry Potter dragged Draco Malfoy 's trousers down past his hips and \mid 0
+-----+
only showing top 10 rows
  text_df= text_df.filter(((text_df.Sentiment =='1') | (text_df.Sentiment =='0')))
  text_df.count()
Out[41]: 6990
  text_df.groupBy("Sentiment").count().show()
|Sentiment|count|
+----+
      0| 3081|
      1| 3909|
+----+
  text_df.printSchema()
|-- Review: string (nullable = true)
|-- Sentiment: string (nullable = true)
```

```
text_df = text_df.withColumn("Label", text_df.Sentiment.cast('float')).drop('Sentiment')
 text_df.orderBy(rand()).show(10,False)
Review
+-----
\midI either LOVE Brokeback Mountain or think it's great that homosexuality \mid1.0 \mid
Da Vinci Code sucks.
|we're gonna like watch Mission Impossible or Hoot.(
                                                              11.0
I love Brokeback Mountain....
                                                             1.0
| {	t I} {	t Would} {	t give you more Kudos on this blog but you had to go and talk abou} | {	t 0.0}
So Brokeback Mountain was really depressing.
                                                            0.0
|No matter how much I love Brokeback Mountain, Crash definitely deserved |1.0
|I think I hate Harry Potter because it outshines much better reading mat|0.0 |
|friday hung out with kelsie and we went and saw The Da Vinci Code SUCKED|0.0\> |
da vinci code sucks...
+----
only showing top 10 rows
 text_df.groupBy('label').count().show()
|label|count|
+----+
1.0| 3909|
| 0.0| 3081|
+----+
 # Add length of the dataframe
 from pyspark.sql.functions import length
 text_df = text_df.withColumn('length', length(text_df['Review']))
 text_df.orderBy(rand()).show(10,False)
+-----
Review
                                                             |Label|length|
|Da Vinci Code = Up, Up, Down, Down, Left, Right, Left, Right, B, A, SUCK | 0.0 | 72
|I love The Da Vinci Code...
                                                             1.0 |27
|Combining the opinion / review from Gary and Gin Zen, The Da Vinci Code |0.0\> |71\>
The Da Vinci Code was absolutely AWESOME!
                                                             1.0 41
                                                                        |You know, the Harry Potter books are decent enough, and I 'm glad the \, |1.0 \, |70
                                                             11.0 | 35
Brokeback mountain was beautiful...
|The Da Vinci Code was absolutely AWESOME!
                                                             11.0 |41
|I| want to be here because I love Harry Potter, and I really want a place |I| 0 |I| 2
I love Brokeback Mountain....
                                                            1.0 |29
                                                                        | \, {\tt "I} \, liked the first "" Mission Impossible."
                                                             1.0 |42
+------
only showing top 10 rows
```

```
text_df.groupBy('Label').agg({'Length': 'mean'}).show()
+----+
|Label| avg(Length)|
1.0 47.61882834484523
0.0|50.95845504706264|
+----+
 # data Cleaning
 tokenization= Tokenizer(inputCol='Review', outputCol='tokens')
 tokenized_df= tokenization.transform(text_df)
 tokenized_df.show()
| Review|Label|length| tokens|
+----+
|The Da Vinci Code...| 1.0| 39|[the, da, vinci, ...|
|this was the firs...| 1.0| 72|[this, was, the, ...|
|i liked the Da Vi...| 1.0| 32|[i, liked, the, d...|
|i liked the Da Vi...| 1.0| 32|[i, liked, the, d...|
|I liked the Da Vi...| 1.0| 72|[i, liked, the, d...|
|that's not even a...| 1.0|
                             72|[that's, not, eve...|
|I loved the Da Vi...| 1.0| 72|[i, loved, the, d...|
|i thought da vinc...| 1.0| 57|[i, thought, da, ...|
|The Da Vinci Code...| 1.0| 45|[the, da, vinci, ...|
|I thought the Da ...| 1.0|
                             51|[i, thought, the,...|
|The Da Vinci Code...| 1.0|
                             68|[the, da, vinci, ...|
|The Da Vinci Code...| 1.0|
                             62|[the, da, vinci, ...|
|then I turn on th...| 1.0|
                             66|[then, i, turn, o...|
|The Da Vinci Code...| 1.0|
                           34|[the, da, vinci, ...|
                             24|[i, love, da, vin...|
|i love da vinci c...| 1.0|
|i loved da vinci ...| 1.0|
                             23|[i, loved, da, vi...|
|TO NIGHT:: THE DA...| 1.0|
                             52|[to, night::, the...|
|THE DA VINCI CODE...| 1.0| 40|[the, da, vinci, ...|
 stopword_removal = StopWordsRemover(inputCol='tokens', outputCol='refined_tokens')
 refined_text_df = stopword_removal.transform(tokenized_df)
 refined_text_df.show()
           Review|Label|length| tokens| refined_tokens|
1
|The Da Vinci Code...| 1.0| 39|[the, da, vinci, ...|[da, vinci, code,...|
|this was the firs...| 1.0| 72|[this, was, the, ...|[first, clive, cu...| | i liked the Da Vi...| 1.0| 32|[i, liked, the, d...|[liked, da, vinci...|
```

```
|i liked the Da Vi...| 1.0|
                             32|[i, liked, the, d...|[liked, da, vinci...|
|I liked the Da Vi...| 1.0|
                             72|[i, liked, the, d...|[liked, da, vinci...|
|that's not even a...| 1.0|
                              72|[that's, not, eve...|[even, exaggerati...|
|I loved the Da Vi...| 1.0|
                              72|[i, loved, the, d...|[loved, da, vinci...|
|i thought da vinc...| 1.0|
                              57|[i, thought, da, ...|[thought, da, vin...|
|The Da Vinci Code...| 1.0|
                              45|[the, da, vinci, ...|[da, vinci, code,...|
|I thought the Da ...| 1.0|
                              51|[i, thought, the,...|[thought, da, vin...|
                              68|[the, da, vinci, \dots|[da, vinci, code,\dots|
|The Da Vinci Code...| 1.0|
                              62|[the, da, vinci, ...|[da, vinci, code,...|
|The Da Vinci Code...| 1.0|
|then I turn on th...| 1.0|
                              66|[then, i, turn, o...|[turn, light, rad...|
|The Da Vinci Code...| 1.0|
                              34|[the, da, vinci, ...|[da, vinci, code,...|
|i love da vinci c...| 1.0|
                              24|[i, love, da, vin...|[love, da, vinci,...|
                              23|[i, loved, da, vi...|[loved, da, vinci...|  
|i loved da vinci ...| 1.0|
|TO NIGHT:: THE DA...| 1.0|
                              52|[to, night::, the...|[night::, da, vin...|
  from pyspark.sql.functions import udf
  from pyspark.sql.types import IntegerType
  from pyspark.sql.functions import *
  len_udf = udf(lambda s: len(s), IntegerType())
  refined_text_df = refined_text_df.withColumn("token_count", len_udf(col('refined_tokens')))
  refined_text_df.orderBy(rand()).show(10)
           Review|Label|length| tokens| refined_tokens|token_count|
|The Da Vinci Code...| 1.0| 30|[the, da, vinci, ...|[da, vinci, code,...|
|i \text{ heard da vinci} \dots| 0.0| 53|[i, \text{ heard, da, vi}\dots|[\text{heard, da, vinci}\dots|
                                                                                 91
|Which is why i sa...| 1.0|
                             72|[which, is, why, ...|[said, silent, hi...|
|I used to hate Ha...| 0.0|
                             28|[i, used, to, hat...|[used, hate, harr...|
|I finished The Da...| 1.0|
                             54|[i, finished, the...|[finished, da, vi...|
|Then snuck into B...| 0.0|
                            72|[then, snuck, int...|[snuck, brokeback...|
|""" I hate Harry ...| 0.0|
                             25|[""", i, hate, ha...|[""", hate, harry...|
                                                                                 41
|Brokeback Mountai...| 0.0|
                              37|[brokeback, mount...|[brokeback, mount...|
                                                                                 3 |
|we're gonna like ...| 1.0|
                              51|[we're, gonna, li...|[gonna, like, wat...|
                                                                                 61
|Brokeback Mountai...| 0.0| 30|[brokeback, mount...|[brokeback, mount...|
                                                                                 3|
+-----
only showing top 10 rows
  count_vec = CountVectorizer(inputCol='refined_tokens', outputCol='features')
  cv_text_df = count_vec.fit(refined_text_df).transform(refined_text_df)
  cv_text_df.select(['refined_tokens', 'token_count', 'features', 'Label']).show(10)
| refined_tokens|token_count| features|Label|
 .-----
|[da, vinci, code,...| 5|(2302,[0,1,4,43,2...| 1.0|
|[first, clive, cu...|
                         9|(2302,[11,51,229,...| 1.0|
5|(2302,[0,1,4,52,3...| 1.0|
|[liked, da, vinci...|
                           5|(2302,[0,1,4,52,3...| 1.0|
|[liked, da, vinci...|
|[liked, da, vinci...|
                           8|(2302,[0,1,4,52,7...| 1.0|
```

```
6|(2302,[46,229,272...| 1.0|
|[even, exaggerati...|
                           8|(2302,[0,1,22,30,...| 1.0|
7|(2302,[0,1,4,228,...| 1.0|
|[loved, da, vinci...|
|[thought, da, vin...|
                        6|(2302,[0,1,4,33,2...| 1.0|
7|(2302,[0,1,4,223,...| 1.0|
|[da, vinci, code,...|
|[thought, da, vin...|
only showing top 10 rows
  # select data for building work
  model_text_df = cv_text_df.select(['features', 'token_count','Label'])
  from pyspark.ml.feature import VectorAssembler
  df_assembler = VectorAssembler(inputCols=['features', 'token_count'], outputCol='features_vec')
  model_text_df =df_assembler.transform(model_text_df)
  model_text_df.printSchema()
root
 |-- features: vector (nullable = true)
 |-- token_count: integer (nullable = true)
 |-- Label: float (nullable = true)
 |-- features_vec: vector (nullable = true)
  from pyspark.ml.classification import LogisticRegression
  # splitting tha train data
  training_df , test_df = model_text_df.randomSplit([0.75, 0.25])
  training_df.groupBy('Label').count().show()
|Label|count|
| 1.0| 2916|
| 0.0| 2296|
+----+
  test_df.groupBy('Label').count().show()
+----+
|Label|count|
| 1.0| 993|
| 0.0| 785|
+----+
```

```
log_reg = LogisticRegression(featuresCol = 'features_vec', labelCol = 'Label').fit(training_df)
```

```
results = log_reg.evaluate(test_df).predictions
```

results.show()

features	token_count L	abel	features_vec	rawPrediction	probability	prediction
(2302,[0,1,4,5,30	5	1.0 (230	3,[0,1,4,5,30 [-	20.033223790660	[1.99379935936333	1.0
(2302,[0,1,4,5,36	5	1.0 (230	3,[0,1,4,5,36 [-	36.496611275475	[1.41163726618194	1.0
(2302,[0,1,4,5,75	5	1.0 (230	3,[0,1,4,5,75 [-	24.189057523923	[3.12482567921014	1.0
(2302,[0,1,4,5,10	6	1.0 (230	3,[0,1,4,5,10 [-	24.600819449628	[2.07014069441782	1.0
(2302,[0,1,4,12,1	10	1.0 (230	3,[0,1,4,12,1 [-	25.485560392607	[8.54597766156175	1.0
(2302,[0,1,4,12,1	5	1.0 (230	3,[0,1,4,12,1 [-	30.350012582238	[6.59412248487400	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302,[0,1,4,12,3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0
(2302, [0, 1, 4, 12, 3	5	1.0 (230	3,[0,1,4,12,3 [-	31.821377509241	[1.51408878001981	1.0

 $from\ pyspark.ml.evaluation\ import\ Binary Classification Evaluator$

```
# confusion matrix
true_positives = results[(results.Label == 1) & (results.prediction == 1)].count()
true_negatives = results[(results.Label == 0) & (results.prediction == 0)].count()
false_positives = results[(results.Label == 0) & (results.prediction == 1)].count()
false_negatives = results[(results.Label == 1) & (results.prediction == 0)].count()
```

```
recall = float(true_positives)/(true_positives + false_negatives)
print(recall)
```

0.9798590130916415

```
precision = float(true_positives)/(true_positives + false_positives)
print(precision)
```

- 0.9778894472361809
- 0.9763779527559056