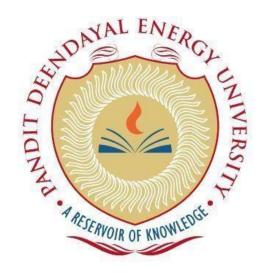
PANDIT DEENDAYAL ENERGY UNIVERSITY SCHOOL OF TECHNOLOGY



Course: Big Data Analytics

Course Code: 20DS506P

LAB MANUAL

M.Tech. (Data Science)

Semester 1

Submitted To:

Dr. Samir patel

Submitted By:

Dev Jethva 23MDS003

Index

Sr. no.	Name of the Lab / Program	Sign
1	Python basics	
2	Basic of Scala	
3	File read	
4	Transformation function	
5	Transformation function example	
6	Transformation function extended	
7	Pyspark join types and data frames	
8	Pyspark and NLP	
9	Linear regression	
10	Page rank calculation	
11	Graph	
12	Structured streaming	
13	Hadoop installation	
14	Pig installation	
15	Kafka installation	

```
(https://databricks.com)
    var num= List(1,2,3,4)
 num: List[Int] = List(1, 2, 3, 4)
   num.head
 res0: Int = 1
   // first two element are show
   num.take(2)
 res3: List[Int] = List(1, 2)
   num.sum
 res4: Int = 10
   var dev = List(1,1,1,2,3,4,2,2)
 dev: List[Int] = List(1, 1, 1, 2, 3, 4, 2, 2)
    // this element show only value to not same like 1,1,2,2,3,4,4 == 1,2,3,4
    dev.distinct
 res7: List[Int] = List(1, 2, 3, 4)
    num(0)
 res8: Int = 1
    num(1) = 10
 command-1298605597657410:1: error: value update is not a member of List[Int]
 num(1)= 10
    dev.size
 res10: Int = 8
    dev.reverse
 res11: List[Int] = List(2, 2, 4, 3, 2, 1, 1, 1)
    dev.min
 res12: Int = 1
```

```
dev.max
res13: Int = 4
  dev.isEmpty
res14: Boolean = false
  var dev1 = List()
  dev1.isEmpty
dev1: List[Nothing] = List()
res16: Boolean = true
  var number = Array(1,2,3,4,5,6,7)
number: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7)
  var lang = Array("Scall", "Python", "Spark")
lang: Array[String] = Array(Scall, Python, Spark)
  lang.head
res17: String = Scall
  lang.tail
res18: Array[String] = Array(Python, Spark)
  number(1)=10
  number
res20: Array[Int] = Array(1, 10, 3, 4, 5, 6, 7)
  import scala.collection.mutable.ArrayBuffer
{\tt import\ scala.collection.mutable.} Array {\tt Buffer}
  var cars = new ArrayBuffer[String]()
cars: scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer()
  cars +="BMW"
```

```
res23: scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer(BMW)
  cars +="Jaguar"
  cars +="Jaguar"
res24: scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer(BMW, Jaguar, Jaguar)
  cars +="TATA"
res 25:\ scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer(BMW,\ Jaguar,\ Jaguar,\ TATA)
  cars.length
res26: Int = 4
  cars.trimEnd(1)
  cars
res28: scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer(BMW, Jaguar, Jaguar)
  cars.insert(2,"Bentley")
  cars
res30: scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer(BMW, Jaguar, Bentley, Jaguar)
  cars.insert(4,"Bugatti")
  cars
res33: scala.collection.mutable.ArrayBuffer[String] = ArrayBuffer(BMW, Jaguar, Bentley, Jaguar, Bugatti)
  val num= List(1,2,3,4,5)
num: List[Int] = List(1, 2, 3, 4, 5)
  num.map(x \Rightarrow x*x)
res34: List[Int] = List(1, 4, 9, 16, 25)
  num.map(y \Rightarrow y*y)
```

```
res35: List[Int] = List(1, 4, 9, 16, 25)
  val a = num.map(x \Rightarrow x+1)
a: List[Int] = List(2, 3, 4, 5, 6)
  // Nested map
  val b = a.map(x \Rightarrow x*x)
b: List[Int] = List(4, 9, 16, 25, 36)
  val c = b.map(x \Rightarrow x-1)
c: List[Int] = List(3, 8, 15, 24, 35)
  val d = c.map(x \Rightarrow -x)
d: List[Int] = List(-3, -8, -15, -24, -35)
  num.map(x \Rightarrow x+1).map(x \Rightarrow x*x).map(x \Rightarrow x-1).map(x \Rightarrow -x)
res36: List[Int] = List(-3, -8, -15, -24, -35)
  val fruits = List("orange", "banana", "Apple")
fruits: List[String] = List(orange, banana, Apple)
  // Length of the words
  fruits.map(x \Rightarrow (x,x.length))
res39: List[(String, Int)] = List((orange,6), (banana,6), (Apple,5))
  \ensuremath{//} how many corrector of the words of their like banana word is geater than 5.
  fruits.filter(x => x.length > 5)
res40: List[String] = List(orange, banana)
  val ratings=List(2.4,5.6,7.8,9.5)
ratings: List[Double] = List(2.4, 5.6, 7.8, 9.5)
  fruits.filter(x => x.length == 5)
res42: List[String] = List(Apple)
  fruits.filter(x => x.length != 5)
```

```
res43: List[String] = List(orange, banana)
  val marks=ratings.map(x \Rightarrow x+10)
marks: List[Double] = List(12.4, 15.6, 17.8, 19.5)
  val marks=ratings.map(x => x*10)
marks: List[Double] = List(24.0, 56.0, 78.0, 95.0)
  val marks1=ratings.filter(x=> x>=60 && x<=74)</pre>
marks1: List[Double] = List()
  // Funtion in Scala
  def add(a:Double = 100 , b:Double = 200) : Double =
    var sum : Double = 0
    sum = a+b
   return sum
  println("Sum:" + add(a = 10, b=20))
Sum:30.0
add: (a: Double, b: Double)Double
  // if value is not given in add funcation return add funcation of a = 100 and b = 200
  println("Sum1:" + add())
Sum1:300.0
   val y = if(x<3)
      println("value of x is less than 3")
    else
      println("value of x is greater than or equal to 3")
value of x is less than 3
x: Int = 1
y: Unit = ()
```

```
var marks = 75
  if(marks>=75){
    println("Distinction")
} else if(marks >= 60 && marks <70){
    println("First class")
} else if(marks >= 50 && marks <60){
    println("Seocnd class")
} else if(marks >= 40 && marks <50){
    println("Pass class")
} else println("Fail")</pre>
```

Distinction
marks: Int = 75

```
// 3 funcations are calls and returns 3 times
def square(x:Double) :Double = {
   return x*x
}
def sumsquare(x:Double, y:Double) :Double = {
   return square(x) + square(x)
}
println("Sum of squares:"+ sumsquare(4,5))
```

Sum of squares:32.0
square: (x: Double)Double
sumsquare: (x: Double, y: Double)Double

```
def time() :Long = {
    println("Inside the funcations")
    return System.nanoTime()
  }
  def exect(t:Long) :Long = {
    println("Inside the Exect funcations")
    println("Time :"+t)
    println("Exiting From exect funcations")
    return t
  }
  println("Main Functions:"+ exect(time()))
```

Inside the funcations
Inside the Exect funcations
Time :4627384344930
Exiting From exect funcations
Main Functions:4627384344930
time: ()Long
exect: (t: Long)Long

```
var i =10
    while(i>0){
    println("hello :"+i)
    i = i-1
}
```

hello :10 hello :9 hello :8 hello :7 hello :6 hello :5 hello :4 hello :3 hello :2

```
hello :1
i: Int = 0
```

```
var a =2
  do{
    println(a)
    a = a +2
}while(a <= 10)</pre>
```

```
;// how to create object and how to create a class

object classEg{
    def main(arg:Array[String]){
        var obj = new NewClass("Hello World")
        obj.sayHi()
    }
}

class NewClass(mssg:String){
    def sayHi()= println(mssg)
}
```

defined object classEg
defined class NewClass

```
for(i <-1 to 10)
println(i)
```

```
// factorial number

var Factorial =1
 var num =5
 while(num>0){
   Factorial = Factorial * num
   num = num-1
 }
 println(Factorial)
```

120 Factorial: Int = 120 num: Int = 0

```
def isPrime(i: Int): Boolean =
    if (i <= 1)
        false
    else if (i == 2)
        true
    else
    !(2 until i).exists(n => i % n == 0)
```

isPrime: (i: Int)Boolean

```
(https://databricks.com)
    // comment in scala
   var a = Array(List(1,2,3), List(1,2,3), List(1,2,3))
   var b = Array(List(4,5,6), List(4,5,6), List(4,5,6))
   var c = Array(Array(0,0,0), Array(0,0,0), Array(0,0,0))
   var sum = 0
   for (i<-0 to 2){
     for (j<-0 to 2){
       sum=0
       for(k<-0\ to\ 2)\{
         sum = sum + (a(i)(k) * b(k)(j))
        c(i)(j)=sum
   }
   println(c)
 a: Array[List[Int]] = Array(List(1, 2, 3), List(1, 2, 3), List(1, 2, 3))
 b: Array[List[Int]] = Array(List(4, 5, 6), List(4, 5, 6), List(4, 5, 6))
 c: Array[Array[Int]] = Array(Array(24, 30, 36), Array(24, 30, 36), Array(24, 30, 36))
 sum: Int = 36
   a(1)(1)
 res2: Int = 2
    var a =Array(Array(1,2,3), List(1,2,3), List(1,2,3))
 a: Array[java.io.Serializable] = Array(Array(1, 2, 3), List(1, 2, 3), List(1, 2, 3))
    a(1)
 res3: java.io.Serializable = List(1, 2, 3)
    a(0)
 res4: java.io.Serializable = Array(1, 2, 3)
Transportation function examples
```

```
import org.apache.spark.rdd.RDD
import org.apache.spark.sql.SparkSession
```

```
import org.apache.spark.rdd.RDD
import org.apache.spark.sql.SparkSession
```

```
object RDDParallelize{
         def main(args: Array[String]): Unit = {
             val\ spark(SparkSession = SparkSession.builder().master("local[1]").appName("SparkByExamples.com").getOrCreate()) appName("SparkByExamples.com").getOrCreate() appName("SparkByExamples.com").getOrCrea
             val rdd:RDD[Int]= spark.sparkContext.parallelize(List(1,2,3,4,5))
             val rddCollect:Array[Int] = rdd.collect()
             println("Number of partitions:" +rdd.getNumPartitions)
              println("Action: First element:" +rdd.first)
             println("Action: RDD converted to Array[Int]: ")
             rddCollect.foreach(println)
     }
defined object RDDParallelize
     val rdda = sc.parallelize(List(1,2,3,4,5))
             val rddb = rdda.collect
             println("Number of Partition:" +rdda.getNumPartitions)
             println("Action: First element:" + rdda.first())
             rdda.foreach(println)
Number of Partition:8
Action: First element:1
rdda: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at command-3160921003574745:1
rddb: Array[Int] = Array(1, 2, 3, 4, 5)
     val rdda = sc.parallelize(List("mumbai", "Delhi", "Chennai", "Kolkatta"))
var rddb = sc.parallelize(Array(1,2,3,4,5,6,7,8,9,10))
rddb: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[2] at parallelize at command-3160921003574747:1
    var rddc = sc.parallelize(Seq.empty[String])
rddc: org.apache.spark.rdd.RDD[String] = ParallelCollectionRDD[3] at parallelize at command-3160921003574748:1
     rdda.collect
res8: Array[String] = Array(mumbai, Delhi, Chennai, Kolkatta)
     rddb.collect
res9: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
     // Collect is action
```

res11: Array[String] = Array()

```
val b = rdda.map(x \Rightarrow (x,1))
  // map, filter is trnasformation
b: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[5] at map at command-3160921003574752:1
  b.collect
res12: Array[(String, Int)] = Array((mumbai,1), (Delhi,1), (Chennai,1), (Kolkatta,1))
  val b = rdda.map((_,1))
b: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[6] at map at command-3160921003574754:1
  b.collect
res13: Array[(String, Int)] = Array((mumbai,1), (Delhi,1), (Chennai,1), (Kolkatta,1))
  val b = rdda.map(x =>(x,x.length))
b: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[7] at map at command-3160921003574756:1
  b.collect
res14: Array[(String, Int)] = Array((mumbai,6), (Delhi,5), (Chennai,7), (Kolkatta,8))
  // Word Length
  val a = sc.parallelize(List(1,2,3,4,5)).map(x=>List(x,x,x)).collect
a: Array[List[Int]] = Array(List(1, 1, 1), List(2, 2, 2), List(3, 3, 3), List(4, 4, 4), List(5, 5, 5))
  val \ a = sc.parallelize(List(1,2,3,4,5)).flatMap(x=>List(x,x,x)).collect
a: Array[Int] = Array(1, 1, 1, 2, 2, 2, 3, 3, 3, 4, 4, 4, 5, 5, 5)
  // What is difference between the map and flatmap
  // applying filter on rdda city data
  val rdda = sc.parallelize(List("Mumbai", "Mumbai", "Delhi", "Chennai", "Kolakatta")).filter(_.equals("Mumbai")).count
rdda: Long = 2
  // applying filter which contain a in city data
  val rdda = sc.parallelize(List("Mumbai", "Delhi", "Chennai", "Kolakatta")).filter(_.contains("e")).collect
rdda: Array[String] = Array(Delhi, Chennai)
```

```
// Creating an rdda with city,Count
       val a = sc.parallelize(List(("Mumbai",4000),( "Delhi", 2000), ("Chennai",1000),("Kolakatta", 7000)))
a: org.apache.spark.rdd.RDD[(String, Int)] = ParallelCollectionRDD[18] \ at parallelize \ at command-3160921003574763:20 \ at parallelize \ at paralleli
       // Perform filter operation where value euals 4000
       // \_.\_1 is for key \_.\_2 is for value
       val a = sc.parallelize(List(("Mumbai",4000),( "Delhi", 2000), ("Chennai",1000),("Kolakatta",
       7000))).filter(_._2.equals(4000)).collect
a: Array[(String, Int)] = Array((Mumbai,4000))
       // Perform filter operation where values grather than 3000
       val \ a = sc.parallelize(List(("Mumbai",4000),( "Delhi", 2000), ("Chennai",1000),("Kolakatta", 7000))).filter(\_._2 > 3000).collect
      // \_.\_1 is for key \_.\_2 is for value
a: Array[(String, Int)] = Array((Mumbai,4000), (Kolakatta,7000))
       // Perform filter which start with C
       val a = sc.parallelize(List(("Mumbai",4000),( "Delhi", 2000), ("Chennai",1000),("Kolakatta", 7000))).filter(a=>
       a._1.startsWith("C")).collect
       // Perform filter by range between 3000, 9000
       val \ b = sc.parallelize(List((4000, "Mumbai"), (2000, "Delhi"), (1000, "Chennai"), (7000, "Kolakatta"))). filter By Range(3000, "Delhi"), (1000, "Chennai"), (1000, "Kolakatta"))). filter By Range(3000, "Delhi"), (1000, "Chennai"), (1000, 
       9000).collect
a: Array[(String, Int)] = Array((Chennai,1000))
b: Array[(Int, String)] = Array((4000, Mumbai), (7000, Kolakatta))
       // sample (flase/true, fraction, seed)
       // false - can not have repeated values
       // true- will have repeated values
       // seed - result will be same if the seed is kept same
       val a = sc.parallelize(1 to 100)
a: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[33] at parallelize at command-3160921003574768:1
       a.sample(false, .2,5).collect
res17: Array[Int] = Array(2, 3, 7, 8, 13, 22, 25, 39, 43, 56, 59, 61, 66, 71, 73, 79, 83, 92)
       a.sample(true, .2,5).collect
res19: Array[Int] = Array(2, 3, 9, 10, 13, 14, 19, 20, 24, 40, 43, 46, 56, 61, 62, 67, 80, 81, 82, 82, 92)
       a.sample(false, .2).collect
```

```
res20: Array[Int] = Array(2, 18, 23, 32, 38, 45, 46, 47, 57, 60, 64, 73, 74, 81, 83, 86, 96)
```

```
a.sample(false,1,5).collect
res21: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29,
30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62,
63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95,
96, 97, 98, 99, 100)
  val a = sc.parallelize(List(1,2,1,1,1,2))
a: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[38] at parallelize at command-3160921003574773:1
  a.sample(true, .4,5).collect
res22: Array[Int] = Array(1)
  val a = sc.parallelize(1 to 7)
a: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[40] at parallelize at command-3160921003574775:1
  val b = sc.parallelize(5 to 10)
b: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[41] at parallelize at command-3160921003574776:1
  a.union(b).collect
res23: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 5, 6, 7, 8, 9, 10)
  \verb"a.intersection(b).collect"
res24: Array[Int] = Array(5, 6, 7)
  a.union(b).distinct.collect
res25: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
  val b = sc.parallelize(1 to 9,3)
b: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[53] at parallelize at command-3160921003574780:1
  b.mapPartitions(x=>List(x.next).iterator).collect
res28: Array[Int] = Array(1, 4, 7)
  val a = sc.parallelize(1 to 9,4)
```

a: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[57] at parallelize at command-3160921003574782:1

5/7

```
a.mapPartitions(x=>List(x.next).iterator).collect
res29: Array[Int] = Array(1, 3, 5, 7)
     def practfunct(index: Int, iter:Iterator[(Int)]): Iterator[String] = {
         iter.toList.map(x => "[index :" +index + ", val : " +x +"]").iterator
practfunct: (index: Int, iter: Iterator[Int])Iterator[String]
     val a = sc.parallelize(List(1,2,3,4,5,6),2)
a.collect
res30: Array[Int] = Array(1, 2, 3, 4, 5, 6)
     a.mapPartitionsWithIndex(practfunct).collect
res 31: Array[String] = Array([index : 0, val : 1], [index : 0, val : 2], [index : 0, val : 3], [index : 1, val : 4], [index : 1, 
5], [index :1, val : 6])
     val a = sc.parallelize(List(1,2,3,4,5,6),3)
a: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[61] at parallelize at command-3160921003574788:1
     a.mapPartitionsWithIndex(practfunct).collect
res32: Array[String] = Array([index :0, val : 1], [index :0, val : 2], [index :1, val : 3], [index :1, val : 4], [index :2, val :
5], [index :2, val : 6])
     val boradcastVar = sc.broadcast(Array(1,2,3))
     boradcastVar.value
boradcastVar:\ org.apache.spark.broadcast.Broadcast[Array[Int]]\ =\ Broadcast(41)
res33: Array[Int] = Array(1, 2, 3)
a: Double = 2500.0
res34: Double = 2500.0
```

♦ databricks2023-08-27 - How to upload File DBFS Example

(https://databricks.com)

Overview

This notebook will show you how to create and query a table or DataFrame that you uploaded to DBFS. DBFS (https://docs.databricks.com/user-guide/dbfs-databricks-file-system.html) is a Databricks File System that allows you to store data for querying inside of Databricks. This notebook assumes that you have a file already inside of DBFS that you would like to read from.

This notebook is written in **Python** so the default cell type is Python. However, you can use different languages by using the <code>%LANGUAGE</code> syntax. Python, Scala, SQL, and R are all supported.

```
# File location and type
file_location = "/FileStore/tables/airline_safety-1.csv"
file_type = "csv"

# CSV options
infer_schema = "false"
first_row_is_header = "false"
delimiter = ","

# The applied options are for CSV files. For other file types, these will be ignored.
df = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)
```

display(df)

Table									
	_c0	_c1	_c2	_c3	_c4	_c5			
1	airline	avail_seat_km_per_week	incidents_85_99	fatal_accidents_85_99	fatalities_85_99	incidents_00_14			
2	Aer Lingus	320906734	2	0	0	0			
3	Aeroflot*	1197672318	76	14	128	6			
4	Aerolineas Argentinas	385803648	6	0	0	1			
5	Aeromexico*	596871813	3	1	64	5			
6	Air Canada	1865253802	2	0	0	2			
7	Air France	3004002661	14	4	79	6			

```
# Create a view or table

temp_table_name = "airline_safety-1_csv"

df.createOrReplaceTempView(temp_table_name)
```

AnalysisException: Invalid view name: airline_safety-1_csv.

```
%sql
/* Query the created temp table in a SQL cell */
select * from `airline_safety-1_csv`
```

(https://databricks.com)

Overview

This notebook will show you how to create and query a table or DataFrame that you uploaded to DBFS. DBFS (https://docs.databricks.com/user-guide/dbfs-databricks-file-system.html) is a Databricks File System that allows you to store data for querying inside of Databricks. This notebook assumes that you have a file already inside of DBFS that you would like to read from.

This notebook is written in **Python** so the default cell type is Python. However, you can use different languages by using the <code>%LANGUAGE</code> syntax. Python, Scala, SQL, and R are all supported.

```
# File location and type
file_location = "/FileStore/tables/file.json"
file_type = "json"
# CSV options
infer_schema = "false"
first_row_is_header = "false"
delimiter = ","
# The applied options are for CSV files. For other file types, these will be ignored.
df = spark.read.format(file_type) \
  .option("inferSchema", infer_schema) \
  .option("header", first_row_is_header) \
  .option("sep", delimiter) \
  .load(file_location)
display(df)
# Create a view or table
temp_table_name = "file_json"
df.createOrReplaceTempView(temp_table_name)
%sql
/* Query the created temp table in a SQL cell */
select * from `file_json`
```

```
(https://databricks.com)
        print("helloworld")
   helloworld
        val accum = sc.longAccumulator("Sum Accumulator")
   accum: org.apache.spark.util.LongAccumulator = LongAccumulator(id: 0, name: Some(Sum Accumulator), value: 0)
   res1: Long = 0
        val broadcastVar = sc.broadcast(Array(1,2,3))
   broadcastVar: \ org.apache.spark.broadcast.Broadcast[Array[Int]] \ = \ Broadcast(\emptyset)
        broadcastVar.value
   res2: Array[Int] = Array(1, 2, 3)
        import\ org. apache. spark. sql. Spark Session
        val spark = SparkSession.builder().appName("SparkExample").master("local").getOrCreate()
    import org.apache.spark.sql.SparkSession
   spark: org.apache.spark.sql.SparkSession = org.apache.spark.sql.SparkSession@eb4526a
        val states = Map(("NY", "New York"),("CA", "California"),("FL", "Florida"))
   states: scala.collection.immutable.Map[String,String] = Map(NY -> New York, CA -> California, FL -> Florida)
        val countries = Map(("USA", "United states of america"), ("IN", "India"))
   countries: scala.collection.immutable.Map[String,String] = Map(USA -> United states of america, IN -> India)
        val broadcaststates = spark.sparkContext.broadcast(states)
   broadcast states: org. apache. spark. broadcast. Broadcast[scala.collection.immutable. Map[String, String]] = Broadcast(1) broadcast states: org. apache. spark. broadcast states: org. apache. org. apa
        val broadcastcountries = spark.sparkContext.broadcast(countries)
   broadcastcountries: org.apache.spark.broadcast.Broadcast[scala.collection.immutable.Map[String,String]] = Broadcast(2)
         val data = Seq(("james", "Suraj", "USA", "CA"),
                                       ("mes", "Sura", "USA", "NY"),
                                       ("megan", "mathews", "USA", "CA"),
                                       ("Maria", "Mohan", "USA", "FL"))
   data: Seq[(String, String, String, String)] = List((james,Suraj,USA,CA), (mes,Sura,USA,NY), (megan,mathews,USA,CA), (Maria,Mohan,US
   A,FL))
        val rdd = spark.sparkContext.parallelize(data)
```

rdd: org.apache.spark.rdd.RDD[(String, String, String, String)] = ParallelCollectionRDD[0] at parallelize at command-260941957792899
1:1

```
val rdd2 = rdd.map(f=> {
  val country = f._3
  val state = f._4
  val fullCountry = broadcastcountries.value.get(country).get
  val fullState = broadcaststates.value.get(state).get
  (f._1, f._2,fullCountry, fullState)
})
```

rdd2: org.apache.spark.rdd.RDD[(String, String, String, String)] = MapPartitionsRDD[1] at map at command-2609419577928992:1

```
println(rdd2.collect().mkString("\n"))
```

```
(james,Suraj,United states of america,California)
(mes,Sura,United states of america,New York)
(megan,mathews,United states of america,California)
(Maria,Mohan,United states of america,Florida)
```

```
val states = Map(("NY", "New York"),("CA", "California"),("FL", "Florida"))
val countries = Map(("USA", "United states of america"), ("IN", "India"))
val broadcaststates = spark.sparkContext.broadcast(states)
val broadcastcountries = spark.sparkContext.broadcast(countries)
val data = Seq(("james", "Suraj", "USA", "CA"),
              ("mes", "Sura", "USA", "NY"),
              ("megan", "mathews", "USA", "CA"),
              ("Maria", "Mohan", "USA", "FL"))
val columns = Seq("First name", "Lastname", "Country", "States")
import spark.sqlContext.implicits._
val df = data.toDF(columns:_*)
val df2 = df.map(row =>{
 val country = row.getString(2)
 val state = row.getString(3)
    val fullcountry = broadcastcountries.value.get(country).get
    val fullstate = broadcaststates.value.get(state).get
    (row.getString(0), row.getString(1), fullcountry,fullstate)
}).toDF(columns: *)
```

```
states: scala.collection.immutable.Map[String,String] = Map(NY -> New York, CA -> California, FL -> Florida)
countries: scala.collection.immutable.Map[String,String] = Map(USA -> United states of america, IN -> India)
broadcaststates: org.apache.spark.broadcast.Broadcast[scala.collection.immutable.Map[String,String]] = Broadcast(4)
broadcastcountries: org.apache.spark.broadcast.Broadcast[scala.collection.immutable.Map[String,String]] = Broadcast(5)
data: Seq[(String, String, String, String)] = List((james,Suraj,USA,CA), (mes,Sura,USA,NY), (megan,mathews,USA,CA), (Maria,Mohan,USA,FL))
columns: Seq[String] = List(First name, Lastname, Country, States)
import spark.sqlContext.implicits._
df: org.apache.spark.sql.DataFrame = [First name: string, Lastname: string ... 2 more fields]
df2: org.apache.spark.sql.DataFrame = [First name: string, Lastname: string ... 2 more fields]
```

```
df2.show(false)
```

```
var longAcc = spark.sparkContext.longAccumulator("SumAccumulator")
longAcc: org.apache.spark.util.LongAccumulator = LongAccumulator(id: 151, name: Some(SumAccumulator), value: 0)
  val rdd = spark.sparkContext.parallelize(Array(1,2,3))
rdd: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[4] at parallelize at command-2609419577928997:1
  rdd.foreach(x \Rightarrow longAcc.add(x))
  println(longAcc.value)
6
  spark.sparkContext.setLogLevel("Error")
  val inputRDD = spark.sparkContext.parallelize(List(("Z", 1),("A", 20),("B", 30),("C", 40),("B", 30),("B", 60)))
inputRDD: org.apache.spark.rdd.RDD[(String, Int)] = ParallelCollectionRDD[5] at parallelize at command-2609419577929001:1
  val listRDD = spark.sparkContext.parallelize(List(1,2,3,4,5,3,2))
listRDD: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[6] at parallelize at command-2609419577929002:1
  // aggregate
  def param0 = (accu:Int, v:Int) \Rightarrow accu + v
  def param1 = (accu1:Int, accu2:Int) => accu1 + accu2
  print("Aggregate :" +listRDD.aggregate(0)(param0, param1))
Aggregate :20param0: (Int, Int) => Int
param1: (Int, Int) => Int
  // aggregate
  def param3 = (accu:Int, v:(String, Int)) => accu + v._2
  def param4 = (accu1:Int, accu2:Int) => accu1 + accu2
  print("Aggregate :" +inputRDD.aggregate(0)(param3, param4))
Aggregate :181param3: (Int, (String, Int)) => Int
param4: (Int, Int) => Int
  // tree Aggregate
  def param8 = (accu:Int, v: Int) \Rightarrow accu + v
  def param9 = (accu1:Int, accu2:Int) => accu1 + accu2
  print("Tree \ Aggregate :" \ +listRDD.treeAggregate(0)(param8, param9))
```

```
Tree Aggregate :20param8: (Int, Int) => Int
param9: (Int, Int) => Int
  // fold action
  println("fold :" +listRDD.fold(0) {(acc,v) =>
  val sum = acc + v
   sum
  })
fold :20
  println("fold :" +inputRDD.fold(("total", 0)) {(acc:(String, Int), v:(String, Int)) =>
  val sum = acc._2+v._2
    ("Total",sum)
  })
fold :(Total,181)
  // what is difference between aggregate, tree aggregate and fold
  val data = listRDD.collect()
data: Array[Int] = Array(1, 2, 3, 4, 5, 3, 2)
  data.foreach(println)
2
  // reduce
  println("Reduce :" +listRDD.reduce(_ + _))
Reduce :20
  println("Reduce :" +listRDD.reduce(_ - _))
Reduce :-10
  \label{eq:println}  \text{println("reduce Alternate :" + inputRDD.reduce((x,y)=> ("Total",x._2+y._2)))} 
reduce Alternate :(Total,181)
  println("tree reduce: " +listRDD.treeReduce(_+_))
```

tree reduce: 20

```
inputRDD.count
res24: Long = 6
  inputRDD.countApproxDistinct()
res26: Long = 5
  listRDD.first
res27: Int = 1
  listRDD.take(2)
res28: Array[Int] = Array(1, 2)
  listRDD.top(2)
res29: Array[Int] = Array(5, 4)
  inputRDD.min._2
res30: Int = 20
  inputRDD.max._2
res31: Int = 1
  inputRDD.min
res32: (String, Int) = (A,20)
  inputRDD.max
res33: (String, Int) = (Z,1)
<console>:8: error: identifier expected but integer literal found.
      partitionied\_rdd = rdd.partitionBy(2, lambda x: x \% 2)
```

```
(https://databricks.com)
   # create spark session
   from pyspark.sql import SparkSession
   spark = SparkSession.builder.appName('nlp').getOrCreate()
   df = spark.createDataFrame([(1, 'I Really Liked this movie'),
                              (2, 'I Would recommend this movie to my friends'),
                             (3, 'movie was alright but acting was horrible'),
                             (4, 'I am never watching that movie ever again')],
                             ['user_id', 'review'])
   df.show(5,False)
 |user_id|review
 |1 | I Really Liked this movie
        |I Would recommend this movie to my friends|
 |2
      |movie was alright but acting was horrible |
 |4 |I am never watching that movie ever again |
   # tokenization
   from pyspark.ml.feature import Tokenizer
   tokenization = Tokenizer(inputCol='review', outputCol='tokens')
   tokenized_df = tokenization.transform(df)
   tokenized_df.show(4,False)
 +----+----+----
     |I Really Liked this movie
                                                 [i, really, liked, this, movie]
 |2
        |I Would recommend this movie to my friends|[i, would, recommend, this, movie, to, my, friends]
 |3
        |movie was alright but acting was horrible |[movie, was, alright, but, acting, was, horrible] |
      \midI am never watching that movie ever again \mid[i, am, never, watching, that, movie, ever, again] \mid
   # Stopwords removal
   from pyspark.ml.feature import StopWordsRemover
```

```
stopword_removal = StopWordsRemover(inputCol='tokens', outputCol='refined_tokens')
  refined_df = stopword_removal.transform(tokenized_df)
  refined_df.select(['user_id', 'tokens', 'refined_tokens']).show(10,False)
|user_id|tokens
                                                              refined tokens
1
       [i, really, liked, this, movie]
                                                             [really, liked, movie]
|2
        [[i, would, recommend, this, movie, to, my, friends] | [recommend, movie, friends]
|3
        |[movie, was, alright, but, acting, was, horrible] |[movie, alright, acting, horrible]|
14
       |[i, am, never, watching, that, movie, ever, again] |[never, watching, movie, ever]
  # Count Vectorizer
  from pyspark.ml.feature import CountVectorizer
  count_vec = CountVectorizer(inputCol='refined_tokens', outputCol='features')
  cv_df= count_vec.fit(refined_df).transform(refined_df)
  cv_df.select(['user_id', 'refined_tokens', 'features']).show(4, False)
|user_id|refined_tokens
                                          features
       |[really, liked, movie] |(11,[0,3,7],[1.0,1.0,1.0])
|[recommend, movie, friends] |(11,[0,6,9],[1.0,1.0,1.0])
12
|3
       |[movie, alright, acting, horrible]|(11,[0,2,4,5],[1.0,1.0,1.0,1.0]) |
4
       |[never, watching, movie, ever] |(11,[0,1,8,10],[1.0,1.0,1.0,1.0])|
  count_vec.fit(refined_df).vocabulary
Out[23]: ['movie',
 'liked'.
 'horrible',
 'friends',
 'ever',
 'alright',
 'recommend',
 'acting',
 'really',
 'never',
 'watching']
```

```
# tf_idf
        from pyspark.ml.feature import HashingTF,IDF
        hashing_vec = HashingTF(inputCol="refined_tokens", outputCol='tf_features')
        hashing_df = hashing_vec.transform(refined_df)
        hashing_df.select(['user_id', 'refined_tokens', 'tf_features']).show(4, False)
|user_id|refined_tokens
                                                                                                                                               tf_features
                           |[really, liked, movie] |(262144,[99172,210223,229264],[1.0,1.0,1.0])
|[recommend, movie, friends] |(262144,[68228,130047,210223],[1.0,1.0,1.0])
 11
 2
                         |[movie, alright, acting, horrible]|(262144,[95685,171118,210223,236263],[1.0,1.0,1.0,1.0])|
|3
4
                      |[never, watching, movie, ever] |(262144,[63139,113673,203802,210223],[1.0,1.0,1.0,1.0])|
        tf_idf_vec = IDF(inputCol='tf_features', outputCol='tf_idf_features')
        tf_idf_df = tf_idf_vec.fit(hashing_df).transform(hashing_df)
        tf_idf_df.select(['user_id', 'tf_idf_features']).show(4, False)
|user_id|tf_idf_features
                           |(262144,[99172,210223,229264],[0.9162907318741551,0.0,0.9162907318741551])
                           \hspace*{0.2in} \hspace*{0
 12
                           |(262144,[95685,171118,210223,236263],[0.9162907318741551,0.9162907318741551,0.0,0.9162907318741551])|
 4
                          |(262144, 63139, 113673, 203802, 210223), [0.9162907318741551, 0.9162907318741551, 0.9162907318741551, 0.0])|
        # classification
        text_df = spark.read.csv('dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/Movie_reviews.csv', inferSchema=True,
        header=True, sep=',')
        text_df.printSchema()
```

```
|-- Review: string (nullable = true)
 |-- Sentiment: string (nullable = true)
  text_df.count()
Out[36]: 7087
  from pyspark.sql.functions import rand
  text_df.orderBy(rand()).show(10,False)
Review
|Harry Potter is AWESOME I don't care if anyone says differently!..
                                                                       1
| \mbox{Harry Potter is AWESOME I don't care if anyone says differently!}..
                                                                        |1
Always knows what I want, not guy crazy, hates Harry Potter..
                                                                       10
|I either LOVE Brokeback Mountain or think it's great that homosexuality|1
|dudeee i LOVED brokeback mountain!!!!
                                                                        |1
|I liked the movie Brokeback Mountain.
                                                                        |1
|I love Harry Potter.
                                                                        |1
li hate brokeback mountain!!
                                                                        10
|Brokeback Mountain was an AWESOME movie.
|Harry Potter dragged Draco Malfoy 's trousers down past his hips and \ |0\>
only showing top 10 rows
  text_df= text_df.filter(((text_df.Sentiment =='1') | (text_df.Sentiment =='0')))
  text_df.count()
Out[41]: 6990
  text_df.groupBy("Sentiment").count().show()
|Sentiment|count|
       0| 3081|
       1| 3909|
  text_df.printSchema()
|-- Review: string (nullable = true)
 |-- Sentiment: string (nullable = true)
```

```
text_df = text_df.withColumn("Label", text_df.Sentiment.cast('float')).drop('Sentiment')
```

text_df.orderBy(rand()).show(10,False)

text_df.groupBy('label').count().show()

only showing top 10 rows

|label|count| +----+ | 1.0| 3909| | 0.0| 3081| +----+

Add length of the dataframe
from pyspark.sql.functions import length

text_df = text_df.withColumn('length', length(text_df['Review']))

text_df.orderBy(rand()).show(10,False)

Review |Label|length| |Da Vinci Code = Up, Up, Down, Down, Left, Right, Left, Right, B, A, SUCK | 0.0 | 72 | I love The Da Vinci Code... 1.0 27 |Combining the opinion / review from Gary and Gin Zen, The Da Vinci Code |0.0> |71 |The Da Vinci Code was absolutely AWESOME! 11.0 41 |You know, the Harry Potter books are decent enough, and I 'm glad the $\,$ |1.0 $\,$ |70 |Brokeback mountain was beautiful... |1.0 |35 |The Da Vinci Code was absolutely AWESOME! 1.0 |41 |I want to be here because I love Harry Potter, and I really want a place|1.0 |72 I love Brokeback Mountain.... 1.0 |29 "I liked the first "" Mission Impossible." 1.0 |42

only showing top 10 rows

```
text_df.groupBy('Label').agg({'Length': 'mean'}).show()
|Label| avg(Length)|
+----+
1.0 47.61882834484523
0.0|50.95845504706264|
 # data Cleaning
  tokenization= Tokenizer(inputCol='Review', outputCol='tokens')
  tokenized_df= tokenization.transform(text_df)
 tokenized_df.show()
+------
     Review|Label|length|
                                tokens
+------
|The Da Vinci Code...| 1.0| 39|[the, da, vinci,.......|
|this was the firs...| 1.0| 72|[this, was, the,......|
|i liked the Da Vi...| 1.0| 32|[i, liked, the, d....|
                          32|[i, liked, the, d....|
72|[i, liked, the, d....|
|i liked the Da Vi...| 1.0| | |
|I liked the Da Vi...| 1.0|
|that's not even a... | 1.0 | 72 | [that's, not, eve. ... |
|I loved the Da Vi...| 1.0| 72|[i, loved, the, d....|
|i thought da vinc...| 1.0| 57|[i, thought, da,......|
|The Da Vinci Code...| 1.0|
                            45|[the, da, vinci,......|
|I thought the Da ...| 1.0|
                            51|[i, thought, the,....|
|The Da Vinci Code...| 1.0| 68|[the, da, vinci,.......|
|The Da Vinci Code...| 1.0|
                          62|[the, da, vinci,......|
|then I turn on th...| 1.0|
                           66|[then, i, turn, o....|
|The Da Vinci Code...| 1.0|
                            34|[the, da, vinci,......|
|i love da vinci c...| 1.0|
                            24|[i, love, da, vin....|
|i loved da vinci ...| 1.0|
                            23|[i, loved, da, vi....|
|TO NIGHT:: THE DA...| 1.0|
                            52|[to, night::, the....|
|THE DA VINCI CODE...| 1.0|
                            40|[the, da, vinci, ...|
 stopword_removal = StopWordsRemover(inputCol='tokens', outputCol='refined_tokens')
 refined_text_df = stopword_removal.transform(tokenized_df)
 refined text df.show()
     Review|Label|length| tokens| refined_tokens|
|The Da Vinci Code...| 1.0| 39|[the, da, vinci, ...|[da, vinci, code,....|
|this was the firs...| 1.0| 72|[this, was, the, ...|[first, clive, cu....|
|i liked the Da Vi...| 1.0| 32|[i, liked, the, d...|[liked, da, vinci....|
```

```
|i liked the Da Vi...| 1.0|
                             32|[i, liked, the, d...|[liked, da, vinci...|
|I liked the Da Vi...| 1.0|
                             72|[i, liked, the, d...|[liked, da, vinci...|
|that's not even a...| 1.0|
                             72|[that's, not, eve...|[even, exaggerati...|
|I loved the Da Vi...| 1.0|
                             72|[i, loved, the, d...|[loved, da, vinci...|
|i thought da vinc...| 1.0|
                             57|[i, thought, da, ...|[thought, da, vin...|
|The Da Vinci Code...| 1.0|
                             45|[the, da, vinci, ...|[da, vinci, code,...|
|I thought the Da ...| 1.0|
                             51|[i, thought, the,...|[thought, da, vin...|
|The Da Vinci Code...| 1.0|
                             68|[the, da, vinci, \ldots|[da, vinci, code,\ldots|
|The Da Vinci Code...| 1.0|
                             62|[the, da, vinci, \dots|[da, vinci, code,\dots|
|then I turn on th...| 1.0|
                             66|[then, i, turn, o...|[turn, light, rad...|
|The Da Vinci Code...| 1.0|
                             34|[the, da, vinci, ...|[da, vinci, code,...|
|i love da vinci c...| 1.0|
                             24|[i, love, da, vin...|[love, da, vinci,...|
|i loved da vinci ...| 1.0|
                             23|[i, loved, da, vi...|[loved, da, vinci...|
|TO NIGHT:: THE DA...| 1.0|
                             52|[to, night::, the...|[night::, da, vin...|
  from pyspark.sql.functions import udf
  from pyspark.sql.types import IntegerType
  from pyspark.sql.functions import *
  len_udf = udf(lambda s: len(s), IntegerType())
  refined_text_df = refined_text_df.withColumn("token_count", len_udf(col('refined_tokens')))
  refined_text_df.orderBy(rand()).show(10)
           Review|Label|length| tokens| refined_tokens|token_count|
|The Da Vinci Code...| 1.0| 30|[the, da, vinci, ...|[da, vinci, code,...|
|i heard da vinci ...| 0.0|
                             53|[i, heard, da, vi...|[heard, da, vinci...|
|Which is why i sa... | 1.0 | 72|[which, is, why, ... | [said, silent, hi... |
| I used to hate Ha... | 0.0 | 28 | [i, used, to, hat... | [used, hate, harr... |
|I finished The Da...| 1.0| 54|[i, finished, the...|[finished, da, vi...|
|Then snuck into B...| 0.0|
                             72|[then, snuck, int...|[snuck, brokeback...|
                                                                                5|
|""" I hate Harry ...| 0.0|
                             25|[""", i, hate, ha...|[""", hate, harry...|
|Brokeback Mountai...| 0.0|
                             37|[brokeback, mount...|[brokeback, mount...|
                                                                                31
|we're gonna like ...| 1.0|
                             51|[we're, gonna, li...|[gonna, like, wat...|
|Brokeback Mountai...| 0.0| 30|[brokeback, mount...|[brokeback, mount...|
only showing top 10 rows
  count_vec = CountVectorizer(inputCol='refined_tokens', outputCol='features')
  cv_text_df = count_vec.fit(refined_text_df).transform(refined_text_df)
  cv_text_df.select(['refined_tokens', 'token_count', 'features', 'Label']).show(10)
refined_tokens|token_count| features|Label|
|[da, vinci, code,...| 5|(2302,[0,1,4,43,2...| 1.0|
                           9|(2302,[11,51,229,...| 1.0|
[first, clive, cu...
|[liked, da, vinci...|
                          5|(2302,[0,1,4,52,3...| 1.0|
                          5|(2302,[0,1,4,52,3...| 1.0|
|[liked, da, vinci...|
                           8|(2302,[0,1,4,52,7...| 1.0|
|[liked, da, vinci...|
```

```
8|(2302,[0,1,22,30,...| 1.0|
7|(2302,[0,1,4.228
|[even, exaggerati...|
|[loved, da, vinci...|
|[thought, da, vin...|
|[da, vinci, code,...|
                                6|(2302,[0,1,4,33,2...| 1.0|
| [da, vinci, code,...| 6|(2302,[0,1,4,33,2...| 1.0|
| [thought, da, vin...| 7|(2302,[0,1,4,223,...| 1.0|
only showing top 10 rows
  \hbox{\tt\# select data for building work}\\
  model_text_df = cv_text_df.select(['features', 'token_count','Label'])
  from pyspark.ml.feature import VectorAssembler
  df_assembler = VectorAssembler(inputCols=['features', 'token_count'], outputCol='features_vec')
  model_text_df =df_assembler.transform(model_text_df)
  model_text_df.printSchema()
root
 |-- features: vector (nullable = true)
 |-- token_count: integer (nullable = true)
 |-- Label: float (nullable = true)
 |-- features_vec: vector (nullable = true)
  from pyspark.ml.classification import LogisticRegression
  # splitting tha train data
  training_df , test_df = model_text_df.randomSplit([0.75, 0.25])
  training_df.groupBy('Label').count().show()
|Label|count|
| 1.0| 2916|
0.0 2296
+---+----+
  test_df.groupBy('Label').count().show()
+---+
|Label|count|
+----+
| 1.0| 993|
| 0.0| 785|
+----+
```

```
log_reg = LogisticRegression(featuresCol = 'features_vec', labelCol = 'Label').fit(training_df)
```

```
results = log_reg.evaluate(test_df).predictions
```

```
results.show()
```

```
features|token_count|Label|
                                               features vec
                                                                 rawPrediction
                                                                                        probability|prediction|
|(2302,[0,1,4,5,30...|
                             5 | 1.0 | (2303, [0,1,4,5,30... | [-20.033223790660... | [1.99379935936333... |
                             5| 1.0|(2303,[0,1,4,5,36...|[-36.496611275475...|[1.41163726618194...|
[(2302,[0,1,4,5,36...]
                                                                                                             1.0
|(2302,[0,1,4,5,75...|
                              5 | 1.0 | (2303, [0,1,4,5,75... | [-24.189057523923... | [3.12482567921014... |
                                                                                                             1.0
|(2302,[0,1,4,5,10...|
                              6 | 1.0 | (2303, [0,1,4,5,10... | [-24.600819449628... | [2.07014069441782... |
                                                                                                             1.0
                              10| 1.0|(2303,[0,1,4,12,1...|[-25.485560392607...|[8.54597766156175...|
[(2302,[0,1,4,12,1...]
                                                                                                             1.0
                             5| 1.0|(2303,[0,1,4,12,1...|[-30.350012582238...|[6.59412248487400...|
|(2302,[0,1,4,12,1...|
                                                                                                             1.0
                               5| 1.0|(2303,[0,1,4,12,3...|[-31.821377509241...|[1.51408878001981...|
|(2302,[0,1,4,12,3...|
                                                                                                             1.0
|(2302,[0,1,4,12,3...|
                               5 | 1.0 | (2303, [0, 1, 4, 12, 3... | [-31.821377509241... | [1.51408878001981... |
                                                                                                             1.0|
|(2302,[0,1,4,12,3...|
                               5 | 1.0 | (2303, [0,1,4,12,3... | [-31.821377509241... | [1.51408878001981... |
                                                                                                             1.0
                               5 | 1.0 | (2303, [0, 1, 4, 12, 3... | [-31.821377509241... | [1.51408878001981... |
|(2302, [0,1,4,12,3...|
                                                                                                             1.0
|(2302,[0,1,4,12,3...|
                               5| 1.0|(2303,[0,1,4,12,3...|[-31.821377509241...|[1.51408878001981...|
                                                                                                             1.0
                               5| 1.0|(2303,[0,1,4,12,3...|[-31.821377509241...|[1.51408878001981...|
|(2302,[0,1,4,12,3...|
                                                                                                             1.0
|(2302,[0,1,4,12,3...|
                               5 | 1.0 | (2303, [0,1,4,12,3... | [-31.821377509241... | [1.51408878001981... |
                                                                                                             1.0
|(2302,[0,1,4,12,3...|
                               5 | 1.0 | (2303, [0,1,4,12,3... | [-31.821377509241... | [1.51408878001981... |
                                                                                                             1.0
                               5 | 1.0 | (2303, [0, 1, 4, 12, 3... | [-31.821377509241... | [1.51408878001981... |
[(2302,[0,1,4,12,3...]
                                                                                                             1.0
|(2302,[0,1,4,12,3...|
                               5| 1.0|(2303,[0,1,4,12,3...|[-31.821377509241...|[1.51408878001981...|
                                                                                                             1.0|
|(2302,[0,1,4,12,3...|
                               5 | 1.0 | (2303, [0,1,4,12,3... | [-31.821377509241... | [1.51408878001981... |
                                                                                                             1.0
|(2302,[0,1,4,12,3...|
                               5 | 1.0 | (2303, [0,1,4,12,3... | [-31.821377509241... | [1.51408878001981... |
                                                                                                             1.0
```

from pyspark.ml.evaluation import BinaryClassificationEvaluator

```
# confusion matrix
true_positives = results[(results.Label == 1) & (results.prediction == 1)].count()
true_negatives = results[(results.Label == 0) & (results.prediction == 0)].count()
false_positives = results[(results.Label == 0) & (results.prediction == 1)].count()
false_negatives = results[(results.Label == 1) & (results.prediction == 0)].count()
```

```
recall = float(true_positives)/(true_positives + false_negatives)
print(recall)
```

0.9798590130916415

```
precision = float(true_positives)/(true_positives + false_positives)
print(precision)
```

- 0.9778894472361809
- 0.9763779527559056

```
(https://databricks.com)
        from pyspark.sql import SparkSession
       spark = SparkSession.builder.appName('lin_reg').getOrCreate()
        # import
       {\tt from~pyspark.ml.regression~import~LinearRegression}
       # Load Dataset
        df = spark.read.format("csv").option("header", "true").load("dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/Linear_regression_d
        # validate the size of the data
       print((df.count(), len(df.columns)))
   (1232, 6)
        # explore the data
       df.printSchema()
      |-- var_1: integer (nullable = true)
      |-- var_2: integer (nullable = true)
      |-- var_3: integer (nullable = true)
      |-- var_4: double (nullable = true)
      |-- var_5: double (nullable = true)
      |-- output: double (nullable = true)
        # view statistical mesures of the data
        df.describe().show(5,False)
   |count |1232 |1232 |1232 |1232
                                                                                                                                                                               1232
                                                                                                                                                                                                                           1232
    | \texttt{mean} | | 715.0819805194806 | 715.0819805194806 | 80.90422077922078 | 0.3263311688311693 | 0.2592727272727275 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.39734172077922014 | 0.3973417207792014 | 0.39734172077922014 | 0.39734172077922014 | 0.39777922014 | 0.39777922014 | 0.397
   |stddev |91.5342940441652 |93.07993263118064|11.458139049993724|0.015012772334166148|0.012907228928000298|0.03326689862173776|
                                       472
    max
                   1009
                                                        1103
                                                                                              116
                                                                                                                                     0.373
                                                                                                                                                                                0.294
        # sneak into the dataset
        df.head(3)
   Out[78]: [Row(var_1=734, var_2=688, var_3=81, var_4=0.328, var_5=0.259, output=0.418),
      Row(var_1=700, var_2=600, var_3=94, var_4=0.32, var_5=0.247, output=0.389),
      Row(var_1=712, var_2=705, var_3=93, var_4=0.311, var_5=0.247, output=0.417)]
        # import corr function from pyspark functions
        from pyspark.sql.functions import corr
        # check for correlation & how to add variable using corr functions
        df.select(corr('var_1', 'output')).show()
```

```
|corr(var_1, output)|
0.9187399607627283
  # import
  from pyspark.ml.linalg import Vector
  from pyspark.ml.feature import VectorAssembler
  # select the columns to create the input vector
  df.columns
Out[82]: ['var_1', 'var_2', 'var_3', 'var_4', 'var_5', 'output']
  # create the vector assembler
  vec_assembler = VectorAssembler(inputCols = ['var_1', 'var_2', 'var_3', 'var_4', 'var_5'], outputCol='features')
  # transform
  features_df = vec_assembler.transform(df)
  features_df.printSchema()
root
|-- var_1: integer (nullable = true)
 |-- var_2: integer (nullable = true)
 |-- var_3: integer (nullable = true)
 |-- var_4: double (nullable = true)
 |-- var_5: double (nullable = true)
 |-- output: double (nullable = true)
 |-- features: vector (nullable = true)
  features_df.select('features').show(5,False)
features
|[734.0,688.0,81.0,0.328,0.259]|
|[700.0,600.0,94.0,0.32,0.247] |
|[712.0,705.0,93.0,0.311,0.247]|
|[734.0,806.0,69.0,0.315,0.26] |
|[613.0,759.0,61.0,0.302,0.24]|
only showing top 5 rows
  model_df = features_df.select('features', 'output')
  model_df.show(5,False)
```

```
print((model_df.count(), len(model_df.columns)))
```

(1232, 2)

Spilt data - Train & Test Data

sec

```
train_df, test_df = model_df.randomSplit([0.7,0.3])
```

```
print((train_df.count(), len(train_df.columns)))
```

(880, 2)

```
print((test_df.count(), len(test_df.columns)))
```

(352, 2)

train_df.describe().show()

Build Linear Regression Model

```
lin_reg = LinearRegression(labelCol='output')
```

```
# fit the linear model on training dataset
  lr_model = lin_reg.fit(train_df)
  lr_model.intercept
Out[100]: 0.18626961441250367
  print(lr_model.coefficients)
[0.0003371220458572444, 5.991498706752434e-05, 0.0002556814872303228, -0.6686443805619972, 0.48044283582971575]
  training_predictions = lr_model.evaluate(train_df)
  {\tt training\_predictions.meanSquaredError}
Out[103]: 0.00013704630021768094
  training_predictions.r2
Out[104]: 0.8772919740056858
  # make predictions on the test data
  test_results= lr_model.evaluate(test_df)
  \ensuremath{\text{\#}}\xspace view the residual errors based on the predictions
  test_results.residuals.show(10)
         residuals
|0.009429260836439135|
|-9.42564848742444...|
|0.013713981224977412|
|-0.01210437809228...|
|-0.00679164860545...|
|0.010048355854590463|
| 8.97412483862492E-4|
-0.012396549858743|
|-0.01181377619641...|
|-0.00117635064512...|
only showing top 10 rows
  # coefficient of determination value for model
  test_results.r2
Out[107]: 0.8476131884871057
```

 ${\tt test_results.rootMeanSquaredError}$

Out[108]: 0.012796007916146983

Out[109]: 0.00016373781859009624

```
(https://databricks.com)
    # Extracting features form the text
   some_text = spark.createDataFrame([
    Apache Spark achieves high performance for both batch
    and streaming data, using a state-of-art DAG scheduler,
    a query optimizer, and a physical execution engine.
   ٠٠٠٠]٫
    Apache Spark is a fast and general-purpose cluseter computing
    system. It provides high-level APIs in Java, Scala, Python
    and R, and an optimized engine that supports general execution
    graphs. It also supports a rich set of higher-level tools including
    Spark SQL for SQL and structured data processing, MLlib for machine
    learning, GraphX for graph processing, and Spark Streaming.
   ٠٠٠٠
   Machine learning is a filed of computer science that often uses
    statistical techniques to give computers the ability to "learn"
    (i.e., progessively improve performance on a specific task)
    with data, without being explicitly programmed.
    ], ['text'])
   import pyspark.ml.feature as feat
    import pyspark.sql.functions as f
    splitter = feat.RegexTokenizer(
        inputCol = 'text',
        outputCol = 'text_split',
        pattern = '\s+|[,.\"]'
    splitter.transform(some_text).select('text_split').take(1)
 Out[7]: [Row(text_split=['apache', 'spark', 'achieves', 'high', 'performance', 'for', 'both', 'batch', 'and', 'streaming', 'data',
 'using', 'a', 'state-of-art', 'dag', 'scheduler', 'a', 'query', 'optimizer', 'and', 'a', 'physical', 'execution', 'engine'])]
    sw_remover = feat.StopWordsRemover(
        inputCol = splitter.getOutputCol()
        ,outputCol = 'no_stopWords'
    )
    sw_remover.transform(splitter.transform(some_text)).select('no_stopWords').take(1)
 Out[9]: [Row(no_stopWords=['apache', 'spark', 'achieves', 'high', 'performance', 'batch', 'streaming', 'data', 'using', 'state-of-ar
 t', 'dag', 'scheduler', 'query', 'optimizer', 'physical', 'execution', 'engine'])]
```

```
hasher = feat.HashingTF(
   inputCol = sw_remover.getOutputCol()
   ,outputCol = 'hashed'
   ,numFeatures = 20
)
```

```
hasher.transform(sw_remover.transform(splitter.transform(some_text))).select('hashed').take(1)
```

Out[11]: [Row(hashed=SparseVector(20, {0: 1.0, 3: 1.0, 6: 1.0, 8: 2.0, 9: 1.0, 11: 1.0, 12: 1.0, 13: 1.0, 15: 2.0, 16: 2.0, 17: 2.0, 18: 1.0, 19: 1.0}))]

```
idf = feat.IDF(
    inputCol= hasher.getOutputCol()
    ,outputCol='features'
)
```

```
idfModel = idf.fit(hasher.transform(sw_remover.transform(splitter.transform(some_text))))
```

```
idf \texttt{Model.transform(hasher.transform(sw\_remover.transform(splitter.transform(some\_text))))}. select(\texttt{'features'}). \\ take(1)
```

Out[14]: [Row(features=SparseVector(20, $\{0: 0.0, 3: 0.0, 6: 0.2877, 8: 0.0, 9: 0.2877, 11: 0.2877, 12: 0.0, 13: 0.0, 15: 0.0, 16: 0.0, 17: 1.3863, 18: 0.6931, 19: 0.0<math>\}))]$

```
from pyspark.ml import Pipeline
```

```
pipeline = Pipeline(stages=[splitter, sw_remover, hasher, idf])
pipelineModel = pipeline.fit(some_text)
```

```
pipelineModel.transform(some_text).select('text', 'features').take(1)
```

Out[17]: [Row(text='\nApache Spark achieves high performance for both batch\nand streaming data, using a state-of-art DAG schedule r,\na query optimizer, and a physical execution engine.\n', features=SparseVector(20, {0: 0.0, 3: 0.0, 6: 0.2877, 8: 0.0, 9: 0.2877, 11: 0.2877, 12: 0.0, 13: 0.0, 15: 0.0, 16: 0.0, 17: 1.3863, 18: 0.6931, 19: 0.0}))]

```
w2v = feat.Word2Vec(
    vectorSize=5
    ,minCount=2
    ,inputCol= sw_remover.getOutputCol()
    ,outputCol='vector'
)
```

```
model = w2v.fit(sw_remover.transform(splitter.transform(some_text)))
```

```
Out[21]: [Row(vector=DenseVector([0.0076, 0.0077, -0.0017, -0.015, -0.004]))]
```

♦ databricks Demo of Page Rank Calculation 2023-10-30 16:15:32

```
(https://databricks.com)
    import org.apache.spark.HashPartitioner
    val links = sc.parallelize(List(("MapR", List("A", "B")), ))
 import org.apache.spark.HashPartitioner
   import org.apache.spark.HashPartitioner
   val links = sc.parallelize(List(("MapR",List("Baidu","Blogger")),("Baidu", List("MapR")),("Blogger",List("Google","Baidu")),("Google",
   var ranks = links.mapValues(v => 1.0)
 import org.apache.spark.HashPartitioner
 links: org.apache.spark.rdd.RDD[(String, List[String])] = ShuffledRDD[1] at partitionBy at command-1660944887008766:3
 ranks: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[2] at mapValues at command-1660944887008766:4
   // val ranks = contributions.reduceByKey((x,y) => x + y).mapValues(v => 0.15+0.85*v)
   val contributions = links.join(ranks).flatMap { case (url, (links, rank)) => links.map(dest => (dest, rank / links.size)) }
 contributions: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[6] at flatMap at command-1660944887008768:1
    val ranks = contributions.reduceByKey((x, y) => x + y).mapValues(v => 0.15 + 0.85*v)
 ranks: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[8] at mapValues at command-1660944887008769:1
    ranks.collect
 res1: Array[(String, Double)] = Array((Google, 0.575), (MapR, 1.8499999999999), (Blogger, 0.575), (Baidu, 1.0))
    val lines= spark.read.textFile("dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/links.txt").rdd
    val iters = 20
    val links = lines.map{ s =>
         val parts = s.split("\\s+")
         (parts(0), parts(1))
        }.distinct().groupByKey().cache()
     var ranks = links.mapValues(v => 1.0)
        for (i <- 1 to iters) {
         val contribs = links.join(ranks).values.flatMap{ case (urls, rank) =>
           val size = urls.size
            urls.map(url => (url, rank / size))
         }
         ranks = contribs.reduceByKey(_ + _).mapValues(0.15 + 0.85 * _)
        val output = ranks.collect()
        output.foreach(tup => println(tup._1 + " has rank:" + tup._2))
        println("======"")
        output.foreach(tup => println(tup. 1 + " has rank:" + f"${tup. 2}%.3f"))
        println("=======")
    ranks.collect()
    val r = ranks.toDF("URL", "PageRank")
    r.show()
```

```
a has rank:1.0
b has rank:1.0455994483347224
c has rank:1.038759531084514
_____
a has rank:1.000
b has rank:1.046
c has rank:1.039
_____
      PageRank|
|URL|
| a| 1.0|
| b|1.0455994483347224|
| c| 1.038759531084514|
lines: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[883] at rdd at command-1660944887008771:1
links: org.apache.spark.rdd.RDD[(String, Iterable[String])] = ShuffledRDD[888] at groupByKey at command-1660944887008771:6
ranks: org.apache.spark.rdd.RDD[(String, Double)] = MapPartitionsRDD[1029] at mapValues at command-1660944887008771:15
output: Array[(String, Double)] = Array((a,1.0), (b,1.0455994483347224), (c,1.038759531084514))
```

```
import org.apache.spark.sql.SparkSession
r.createOrReplaceTempView("Table_2")
val r1=sqlContext.sql("select PageRank from Table_2 where PageRank <2")</pre>
```

```
import org.apache.spark.sql.SparkSession
r1: org.apache.spark.sql.DataFrame = [PageRank: double]
formattedArray: org.apache.spark.sql.Dataset[String] = [value: string]
res16: Array[String] = Array(1.000, 1.046, 1.039)
```

⊗ databricksGraph Example 2023-10-23 11:00:49

```
(https://databricks.com)
    %python
    # file location and type
    file_location = "dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/vertex.csv"
    file_type = "csv"
    # csv Options
    infer_schema = "false"
    first_row_is_header = "false"
    delimiter = ","
    # the applied options are for csv files.for other file types, these will be ignored.
    df = spark.read.format(file_type) \
        .option("inferSchema", infer_schema)\
        .option("header",first_row_is_header)\
        .option("sep",delimiter)\
        .load(file_location)
    display(df)
```

Table	•				
	_c0	_	_c1	_c2	
1	6		1	Sister	
2	1		2	Husband	
3	2		1	Wife	
4	5		1	Daughter	
5	5		2	Daughter	

6	3	1	Son
7	3	2	Son
12 rows	WS		

```
%scala
import org.apache.spark.rdd.RDD
```

import org.apache.spark.rdd.RDD

```
%scala
import org.apache.spark.graphx._
```

import org.apache.spark.graphx._

```
%scala
val vertexRDD = sc.textFile("dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/vertex.csv")
val edgeRDD = sc.textFile("dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/edges.csv")
edgeRDD.collect()
```

vertexRDD: org.apache.spark.rdd.RDD[String] = dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/vertex.csv MapPartitionsRDD[23]
at textFile at command-223217164383677:1
edgeRDD: org.apache.spark.rdd.RDD[String] = dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/edges.csv MapPartitionsRDD[25] at
textFile at command-223217164383677:2
res0: Array[String] = Array(6,1,Sister, 1,2,Husband, 2,1,Wife, 5,1,Daughter, 5,2,Daughter, 3,1,Son, 3,2,Son, 4,1,Friend, 1,5,Father,
1,3,Father, 2,5,Mother, 2,3,Mother)

```
%scala
vertexRDD.collect()
```

res1: Array[String] = Array(1,Jacob,48, 2,Jessica,45, 3,Andrew,25, 4,Ryan,53, 5,Emily,22, 6,Lily,52)

vertices: org.apache.spark.rdd.RDD[(org.apache.spark.graphx.VertexId, (String, String))] = MapPartitionsRDD[26] at map at command-22
3217164383679:1
res2: Array[(org.apache.spark.graphx.VertexId, (String, String))] = Array((1,(Jacob,48)), (2,(Jessica,45)), (3,(Andrew,25)), (4,(Ryan,53)), (5,(Emily,22)), (6,(Lily,52)))

edges: org.apache.spark.rdd.RDD[org.apache.spark.graphx.Edge[String]] = MapPartitionsRDD[27] at map at command-223217164383680:1 res3: Array[org.apache.spark.graphx.Edge[String]] = Array(Edge(6,1,Sister), Edge(1,2,Husband), Edge(2,1,Wife), Edge(5,1,Daughter), Edge(5,2,Daughter), Edge(3,1,Son), Edge(3,2,Son), Edge(4,1,Friend), Edge(1,5,Father), Edge(1,3,Father), Edge(2,5,Mother), Edge(2,3,Mother))

```
%scala

val default = ("unknown", "missing")

val graph = Graph(vertices, edges, default)
```

```
default: (String, String) = (unknown,missing)
graph: org.apache.spark.graphx.Graph[(String, String),String] = org.apache.spark.graphx.impl.GraphImpl@27608e58
```

defined class MoviesWatched

movies: org.apache.spark.rdd.RDD[(org.apache.spark.graphx.VertexId, MoviesWatched)] = ParallelCollectionRDD[40] at parallelize at co mmand-223217164383682:3

```
%scala
val movieOuterJoinedGraph = graph.outerJoinVertices(movies)((_,name, movies) => (name, movies))
```

movieOuterJoinedGraph: org.apache.spark.graphx.Graph[((String, String), Option[MoviesWatched]),String] = org.apache.spark.graphx.imp l.GraphImpl@42a47cbc

```
%scala
movieOuterJoinedGraph.vertices.map(t =>t).collect.foreach(println)
```

```
(4,((Ryan,53),None))
(6,((Lily,52),None))
(2,((Jessica,45),Some(MoviesWatched(Titanic,Love))))
(1,((Jacob,48),Some(MoviesWatched(Toy Story 3,Kids))))
(3,((Andrew,25),Some(MoviesWatched(The Hangover ,Comedy))))
(5,((Emily,22),None))
```

```
%scala
val movieOuterJoinedGraph = graph.outerJoinVertices(movies)((_,name, movies) => (name, movies.getOrElse(MoviesWatched("NA",
"NA"))))
```

movieOuterJoinedGraph: org.apache.spark.graphx.Graph[((String, String), MoviesWatched),String] = org.apache.spark.graphx.impl.GraphI
mpl@7980553c

```
%scala
movieOuterJoinedGraph.vertices.map(t=>t).collect.foreach(println)
```

```
(4,((Ryan,53),MoviesWatched(NA,NA)))
(6,((Lily,52),MoviesWatched(NA,NA)))
(2,((Jessica,45),MoviesWatched(Titanic,Love)))
(1,((Jacob,48),MoviesWatched(Toy Story 3,Kids)))
(3,((Andrew,25),MoviesWatched(The Hangover ,Comedy)))
(5,((Emily,22),MoviesWatched(NA,NA)))
```

```
%scala
val tCount = graph.triangleCount().vertices
```

tCount: org.apache.spark.graphx.VertexRDD[Int] = VertexRDDImpl[102] at RDD at VertexRDD.scala:57

```
println(tCount.collect().mkString("\n"))
(4,0)
(6,0)
(2,2)
(1,2)
(3,1)
(5,1)
     %scala
     val iterations = 1000
iterations: Int = 1000
     val connected = graph.connectedComponents().vertices
connected: org.apache.spark.graphx.VertexRDD[org.apache.spark.graphx.VertexId] = VertexRDDImpl[126] at RDD at VertexRDD.scala:57
     %scala
     val connections = graph.stronglyConnectedComponents(iterations).vertices
connections: org.apache.spark.graphx.VertexRDD[org.apache.spark.graphx.VertexId] = VertexRDDImpl[377] at RDD at VertexRDD.scala:57
    %scala
     val connByPerson = vertices.join(connected).map{ case(id, ((person,age), conn)) => (conn, id,person)}
     val connByPersonS = vertices.join(connected).map{ case(id, ((person,age), conn)) => (conn, id,person)}
     {\tt connByPerson.collect().foreach\{\ case\ (conn,\ id,\ person) => println(f"Weak\ $conn\ $id\ $person")}\}
Weak 1 4 Ryan
Weak 1 6 Lily
Weak 1 2 Jessica
Weak 1 1 Jacob
Weak 1 3 Andrew
Weak 1 5 Emily
connByPerson: org.apache.spark.rdd.RDD[(org.apache.spark.graphx.VertexId, org.apache.spark.graphx.VertexId, String)] = MapPartitions (Application org.apache.spark.graphx.vertexId, String) = MapPartition org.apache.spark.graphx.vertexId
RDD[427] at map at command-223217164383692:1
connByPersonS: org.apache.spark.rdd.RDD[(org.apache.spark.graphx.VertexId, org.apache.spark.graphx.VertexId, String)] = MapPartition
sRDD[431] at map at command-223217164383692:2
     println("Vertices count:" +graph.vertices.count)
Vertices count:6
     %scala
     println("Vertices count:" +graph.edges.count)
Vertices count:12
     %scala
     val cnt1 = graph.vertices.filter{ case (id,(name, age)) => age.toLong > 40}.count
```

```
cnt1: Long = 4
```

```
%scala
val cnt2 = graph.edges.filter{ case Edge(from ,to, property) => property == "Father" | property == "Mother"}.count
```

cnt2: Long = 4

```
%scala
def max (a: (VertexId, Int), b: (VertexId, Int)) : (VertexId, Int)={
   if(a._2 > b._2) a else b
}
```

max: (a: (org.apache.spark.graphx.VertexId, Int), b: (org.apache.spark.graphx.VertexId, Int))(org.apache.spark.graphx.VertexId, Int)

```
%scala
val maxInDegree : (VertexId, Int) = graph.inDegrees.reduce(max)
val maxOutDegree : (VertexId, Int) = graph.outDegrees.reduce(max)
val maxDegree : (VertexId, Int) = graph.degrees.reduce(max)
```

```
maxInDegree: (org.apache.spark.graphx.VertexId, Int) = (1,5)
maxOutDegree: (org.apache.spark.graphx.VertexId, Int) = (2,3)
maxDegree: (org.apache.spark.graphx.VertexId, Int) = (1,8)
```

```
%scala
val minDegrees = graph.outDegrees.filter(_._2 <=1)
minDegrees.collect()</pre>
```

minDegrees: org.apache.spark.graphx.VertexRDD[Int] = VertexRDDImpl[451] at RDD at VertexRDD.scala:57 res10: Array[(org.apache.spark.graphx.VertexId, Int)] = Array((4,1), (6,1))

```
%scala
graph.triplets.map(
  triplet => triplet.srcAttr._1 + "is the " + triplet.attr + "of " + triplet.dstAttr._1
).collect.foreach(println)
```

Jacobis the Husbandof Jessica
Jessicais the Wifeof Jacob
Andrewis the Sonof Jacob
Emilyis the Daughterof Jacob
Emilyis the Daughterof Jessica
Lilyis the Sisterof Jacob
Jacobis the Fatherof Andrew
Jacobis the Fatherof Emily
Jessicais the Motherof Emily
Andrewis the Sonof Jessica
Ryanis the Friendof Jacob

```
%scala
print(sc.version)
```

3.3.2

```
Python interpreter will be restarted.

Collecting networkx

Downloading networkx-3.2-py3-none-any.whl (1.6 MB)

Installing collected packages: networkx

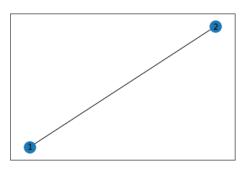
Successfully installed networkx-3.2

Python interpreter will be restarted.
```

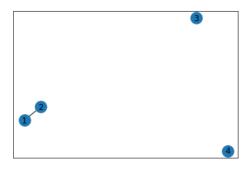
```
import networkx as nx
```

```
# create an Empty undirected graph
G = nx.Graph()

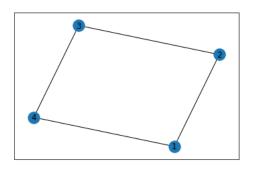
import matplotlib.pyplot as plt
# adding
G.add_edge(1,2)
nx.draw_networkx(G)
plt.show()
```



```
G.add_nodes_from([3,4])
nx.draw_networkx(G)
plt.show()
```



```
G.add_edge(3,4)
G.add_edges_from([(2,3), (4,1)])
nx.draw_networkx(G)
plt.show()
```



G.nodes

Out[5]: NodeView((1, 2, 3, 4))

G.edges

Out[6]: EdgeView([(1, 2), (1, 4), (2, 3), (3, 4)])

list(nx.generate_adjlist(G))

Out[7]: ['1 2 4', '2 3', '3 4', '4']

nx.to_dict_of_lists(G)

Out[8]: {1: [2, 4], 2: [1, 3], 3: [4, 2], 4: [3, 1]}

 $nx.to_edgelist(G)$

 ${\tt Out[10]: EdgeDataView([(1, 2, {\}}), (1, 4, {\}}), (2, 3, {\}}), (3, 4, {\}})])}\\$

nx.to_pandas_adjacency(G)

```
        1
        2
        3
        4

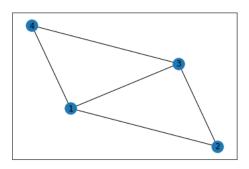
        1
        0.0
        1.0
        0.0
        1.0

        2
        1.0
        0.0
        1.0
        0.0

        3
        0.0
        1.0
        0.0
        1.0

        4
        1.0
        0.0
        1.0
        0.0
```

```
G.add_edge(1,3)
nx.draw_networkx(G)
plt.show()
```



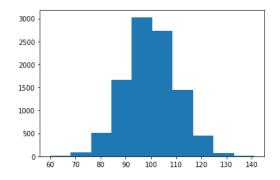
```
# print(nx.to_scipy_sparse_matrix(G))
```

G.degree

Out[13]: DegreeView({1: 3, 2: 2, 3: 3, 4: 2})

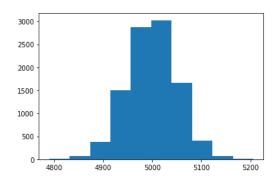
```
k = nx.fast_gnp_random_graph(10000, 0.01).degree()
plt.hist(list(dict(k).values()))
```

```
Out[14]: (array([ 5., 82., 502., 1671., 3028., 2730., 1450., 455., 71., 6.]),
array([ 60., 68.1, 76.2, 84.3, 92.4, 100.5, 108.6, 116.7, 124.8, 132.9, 141. ]),
<BarContainer object of 10 artists>)
```

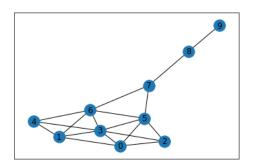


```
k = nx.fast_gnp_random_graph(10000, 0.50).degree()
plt.hist(list(dict(k).values()))
```

```
Out[15]: (array([ 6., 64., 379., 1507., 2879., 3025., 1660., 407., 65. 8.]),
array([4790., 4831.6, 4873.2, 4914.8, 4956.4, 4998., 5039.6, 5081.2, 5122.8, 5164.4, 5206.]),
<BarContainer object of 10 artists>)
```



```
G = nx.krackhardt_kite_graph()
nx.draw_networkx(G)
plt.show()
```

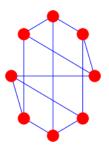


```
print(nx.has_path(G, source =1 , target=9))
print(nx.shortest_path(G, source =1 , target=9))
print(nx.shortest_path_length(G, source =1 , target=9))
print(list(nx.shortest_simple_paths(G, source =1 , target=9)))
paths = list(nx.all_pairs_shortest_path(G))
paths[5][1]
```

```
True
[1, 6, 7, 8, 9]
[[1, 6, 7, 8, 9], [1, 0, 5, 7, 8, 9], [1, 6, 5, 7, 8, 9], [1, 3, 5, 7, 8, 9], [1, 4, 6, 7, 8, 9], [1, 3, 6, 7, 8, 9], [1, 0, 2,
5, 7, 8, 9], [1, 0, 5, 6, 7, 8, 9], [1, 6, 3, 5, 7, 8, 9], [1, 3, 5, 6, 7, 8, 9], [1, 4, 3, 5, 7, 8, 9], [1, 4, 6, 5, 7, 8, 9],
[1, 3, 0, 5, 7, 8, 9], [1, 3, 6, 5, 7, 8, 9], [1, 0, 3, 5, 7, 8, 9], [1, 4, 3, 6, 7, 8, 9], [1, 3, 2, 5, 7, 8, 9], [1, 0, 3, 6, 9], [1, 0, 1, 0]
7, 8, 9], [1, 3, 4, 6, 7, 8, 9], [1, 0, 2, 3, 5, 7, 8, 9], [1, 0, 2, 5, 6, 7, 8, 9], [1, 0, 5, 3, 6, 7, 8, 9], [1, 6, 4, 3, 5, 7,
8, 9], [1, 6, 3, 0, 5, 7, 8, 9], [1, 4, 3, 5, 6, 7, 8, 9], [1, 4, 6, 3, 5, 7, 8, 9], [1, 3, 0, 2, 5, 7, 8, 9], [1, 3, 0, 5, 6, 7,
8, \ 9], \ [1, \ 0, \ 3, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 4, \ 3, \ 0, \ 5, \ 7, \ 8, \ 9], \ [1, \ 4, \ 3, \ 6, \ 5, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 0, \ 5, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 5, \ 6, \ 7, \ 8, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 2, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9], \ [1, \ 3, \ 3, \ 9
8, 9], [1, 0, 3, 2, 5, 7, 8, 9], [1, 0, 3, 6, 5, 7, 8, 9], [1, 3, 4, 6, 5, 7, 8, 9], [1, 0, 2, 3, 6, 7, 8, 9], [1, 6, 3, 2, 5, 7,
8, 9], [1, 4, 3, 2, 5, 7, 8, 9], [1, 0, 3, 4, 6, 7, 8, 9], [1, 0, 2, 3, 5, 6, 7, 8, 9], [1, 0, 2, 5, 3, 6, 7, 8, 9], [1, 0, 5, 2,
3, 6, 7, 8, 9], [1, 0, 5, 3, 4, 6, 7, 8, 9], [1, 6, 4, 3, 0, 5, 7, 8, 9], [1, 6, 3, 0, 2, 5, 7, 8, 9], [1, 4, 6, 3, 0, 5, 7, 8,
9], [1, 3, 0, 2, 5, 6, 7, 8, 9], [1, 4, 3, 0, 2, 5, 7, 8, 9], [1, 4, 3, 0, 5, 6, 7, 8, 9], [1, 3, 2, 0, 5, 6, 7, 8, 9], [1, 0, 3,
2, 5, 6, 7, 8, 9], [1, 0, 2, 3, 4, 6, 7, 8, 9], [1, 0, 2, 3, 6, 5, 7, 8, 9], [1, 6, 3, 2, 0, 5, 7, 8, 9], [1, 4, 3, 2, 0, 5, 7,
8, 9], [1, 4, 3, 2, 5, 6, 7, 8, 9], [1, 0, 3, 4, 6, 5, 7, 8, 9], [1, 6, 4, 3, 2, 5, 7, 8, 9], [1, 4, 6, 3, 2, 5, 7, 8, 9], [1, 0,
2, 5, 3, 4, 6, 7, 8, 9], [1, 0, 5, 2, 3, 4, 6, 7, 8, 9], [1, 6, 4, 3, 0, 2, 5, 7, 8, 9], [1, 4, 6, 3, 0, 2, 5, 7, 8, 9], [1, 4,
3, 0, 2, 5, 6, 7, 8, 9], [1, 0, 2, 3, 4, 6, 5, 7, 8, 9], [1, 4, 3, 2, 0, 5, 6, 7, 8, 9], [1, 6, 4, 3, 2, 0, 5, 7, 8, 9], [1, 4,
6, 3, 2, 0, 5, 7, 8, 9]]
Out[18]: {5: [5],
  0: [5, 0],
  2: [5, 2],
```

```
# importance of nodes inside the network
  nx.betweenness_centrality(G)
Out[19]: {0: 0.023148148148148143,
1: 0.023148148148148143,
2: 0.0,
3: 0.10185185185185183,
4: 0.0,
5: 0.23148148148148148,
6: 0.23148148148148148.
7: 0.3888888888888884,
8: 0.2222222222222,
9: 0.0}
  nx.degree_centrality(G)
2: 0.33333333333333333333,
5: 0.55555555555556,
6: 0.555555555555556,
7: 0.3333333333333333333
8: 0.2222222222222,
9: 0.11111111111111111111
  nx.closeness_centrality(G)
Out[21]: {0: 0.5294117647058824,
1: 0.5294117647058824,
2: 0.5,
3: 0.6,
4: 0.5,
5: 0.6428571428571429,
6: 0.6428571428571429,
7: 0.6,
8: 0.42857142857142855,
9: 0.3103448275862069}
  nx.harmonic_centrality(G)
1: 6.083333333333333,
2: 5.583333333333333,
3: 7.083333333333333,
4: 5.583333333333333,
5: 6.833333333333333,
6: 6.83333333333333,
7: 6.0,
9: 3.416666666666665}
  nx.eigenvector\_centrality(G)
Out[23]: {0: 0.3522089813920359,
1: 0.3522089813920358,
2: 0.28583473531632403,
3: 0.48102048812210046,
4: 0.28583473531632403,
5: 0.3976910106255469,
```

- 6: 0.39769101062554685, 7: 0.19586185175360382,
- 8: 0.048074775014202924,
- 9: 0.011164058575824235}



```
(https://databricks.com)
    %python
    pip install networkx
 Python interpreter will be restarted.
 Collecting networkx
   Downloading networkx-3.2-py3-none-any.whl (1.6 MB)
 Installing collected packages: networkx
 Successfully installed network x-3.2
 Python interpreter will be restarted.
    from pylab import rcParams
   rcParams['figure.figsize']= (3,3)
    def nice_print(v , digits = 3):
       format = '%%.%df' %digits
        print(', '.join([format % e for e in v]))
    nice_print([.12333122, .13432221, .64442143])
    nice_print([.12333122, .13432221, .64442143], digits = 4)
 0.123, 0.134, 0.644
 0.1233, 0.1343, 0.6444
    labels = ['A',
              'B',
              'C',
               'D',
              'E',
              'F',
              'G']
    pages = range(len(labels))
    positions = [(0,1),
                 (0,2),
                 (2,2),
                 (0,0),
```

```
Out[10]: {0: 'A', 1: 'B', 2: 'C', 3: 'D', 4: 'E', 5: 'F', 6: 'G'}
```

this dictionary accosciates the numbers in pages to labels
page_labels = {p: 1 for p, 1 in zip(pages, labels)}

(1,0), (2,0), (1,1)]

page_labels

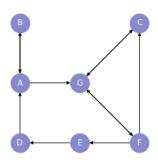
```
import networkx as nx
import matplotlib.pyplot as plt

g = nx.DiGraph()

for p in pages:
    node = g.add_node(p)

for (a,b) in links:
    g.add_edge(pages[a], pages[b])
```

```
plt.clf()
display(nx.draw(g, with_labels = True, labels = page_labels, node_size = 800, node_color = '#8888CC', font_color = 'white', pos
= positions))
```



```
adjacency ={}
for u in range(len(pages)):
    adjacency[u] = []

for (a,b) in links:
    adjacency[a].append(b)

print(adjacency)
```

```
{0: [1, 6], 1: [0], 2: [6], 3: [0], 4: [3], 5: [2, 6, 4], 6: [2, 5]}
```

```
connection_matrix =[]
for a in adjacency:
   for b in adjacency[a]:
        connection_matrix.append((b,a,1./len(adjacency[a])))
connection_matrix
```

```
Out[17]: [(1, 0, 0.5),

(6, 0, 0.5),

(0, 1, 1.0),

(6, 2, 1.0),

(0, 3, 1.0),

(3, 4, 1.0),

(2, 5, 0.3333333333333333),

(4, 5, 0.333333333333333),

(4, 5, 0.333333333333333),

(5, 6, 0.5),

(5, 6, 0.5)]
```

```
links_RDD = sc.parallelize(connection_matrix).cache()
```

```
links_RDD.take(3)
```

```
Out[19]: [(1, 0, 0.5), (6, 0, 0.5), (0, 1, 1.0)]
```

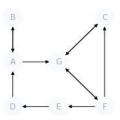
```
import numpy as np
n = len(pages)
page_rank = np.ones(n)/n
old_page_rank = np.ones(n)
print("Page rank is", page_rank)
print("Old Page rank is", old_page_rank)
```

```
Page rank is [0.14285714\ 0.14285714\ 0.14285714\ 0.14285714\ 0.14285714\ 0.14285714] Old Page rank is [1.\ 1.\ 1.\ 1.\ 1.\ 1.\ 1.]
```

```
def l2distance(v,q):
    if len(v) != len(q):
        raise ValueError('Cannot compute the distance of two vectors of different size')
    return sum([(q_el-v_el)**2 for v_el, q_el in zip(v,q)])
```

```
tolerance = 10e-7
max_iterations = 1000
iteration = 0
print("iiiiiii")
while(12distance(old_page_rank, page_rank) >= tolerance and iteration < max_iterations):
    old_page_rank = page_rank
    page_rank_values = links_RDD.map(lambda x:(x[0], x[2]*page_rank[x[1]])).reduceByKey(lambda a, b: a+b).sortByKey().collect()
    page_rank = np.array([c for (i,c) in page_rank_values])
    nice_print(page_rank)
    print("Page rank", page_rank)
    iteration += 1</pre>
```

iiiiiiii



```
(https://databricks.com)
    PySpark Join is used to combine two DataFrames and by chaining these you can join multiple DataFrames; it supports all basic join type
    OUTER, RIGHT OUTER, LEFT ANTI, LEFT SEMI, CROSS, SELF JOIN. PySpark Joins are wider transformations that involve data shuffling across
    PySpark SQL Joins comes with more optimization by default (thanks to DataFrames) however still there would be some performance issues
    In this PySpark SQL Join section, you will learn different Join syntaxes and using different Join types on two or more DataFrames and
    PySpark Join Syntax
   PySpark Join Types
    Inner Join DataFrame
   Full Outer Join DataFrame
    Left Outer Join DataFrame
    Right Outer Join DataFrame
   Left Anti Join DataFrame
    Left Semi Join DataFrame
    Self Join DataFrame
   Using SQL Expression
    1. PySpark Join Syntax
    PySpark SQL join has a below syntax and it can be accessed directly from DataFrame.
    join() operation takes parameters as below and returns DataFrame.
    join(self, other, on=None, how=None)
    param other: Right side of the join
    param on: a string for the join column name
    param how: default inner. Must be one of inner, cross, outer,full, full_outer, left, left_outer, right, right_outer,left_semi, and lef
    You can also write Join expression by adding where() and filter() methods on DataFrame and can have Join on multiple columns.
    2. PvSpark Join Types
    Below are the different Join Types PySpark supports.
    Join String
                                         Equivalent SQL Join
    inner
                                       INNER JOIN
   outer, full, fullouter, full_outer FULL OUTER JOIN
   left, leftouter, left_outer
                                        LEFT JOIN
                                         RIGHT JOIN
   right, rightouter, right_outer
    cross
    anti, leftanti, left anti
    semi, leftsemi, left_semi
    Before we jump into PySpark SQL Join examples, first, let's create an "emp" and "dept" DataFrames. here, column "emp_id" is unique on
    emp_dept_id from emp has a reference to dept_id on dept dataset.
    emp = [(1,"Smith",-1,"2018","10","M",3000), \
        (2,"Rose",1,"2010","20","M",4000), \
        (3,"Williams",1,"2010","10","M",1000), \
        (4,"Jones",2,"2005","10","F",2000), \
        (5,"Brown",2,"2010","40","",-1), \
          (6,"Brown",2,"2010","50","",-1) \
    empColumns = ["emp_id","name","superior_emp_id","year_joined", \
           "emp_dept_id", "gender", "salary"]
    empDF = spark.createDataFrame(data=emp, schema = empColumns)
    empDF.printSchema()
    empDF.show(truncate=False)
    dept = [("Finance",10), \
        ("Marketing",20), \
        ("Sales",30), \
```

("IT",40) \

```
deptColumns = ["dept_name","dept_id"]
 deptDF = spark.createDataFrame(data=dept, schema = deptColumns)
 deptDF.printSchema()
 deptDF.show(truncate=False)
 '''This prints "emp" and "dept" DataFrame to the console. Refer complete example below on how to create spark object.'''
root
|-- emp_id: long (nullable = true)
|-- name: string (nullable = true)
|-- superior_emp_id: long (nullable = true)
|-- year_joined: string (nullable = true)
|-- emp_dept_id: string (nullable = true)
|-- gender: string (nullable = true)
|-- salary: long (nullable = true)
+----+
|emp_id|name |superior_emp_id|year_joined|emp_dept_id|gender|salary|
2018
                                |10
    |Smith |-1
                                          M |3000 |
11
                   2010
                                20
                                              4000
2
     Rose |1
     |Williams|1
                                |10
                                         |M |1000 |
13
                      2005
    |Jones |2
|Brown |2
                                 10
                                          |F
                                               |2000 |
14
|5
                       2010
                                 40
                                                |-1
    |Brown |2
                       2010
                                |50
                                                |-1
6
                                          root
```

3. PySpark Inner Join DataFrame
Inner join is the default join in PySpark and it's mostly used. This joins two datasets on key columns, where keys don't match
the rows get dropped from both datasets (emp & dept).

When we apply Inner join on our datasets, It drops "emp_dept_id" 50 from "emp" and "dept_id" 30 from "dept" datasets. Below is
the result of the above Join expression.
...
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"inner") \
.show(truncate=False)

	•	•	emp_id year_joine	d emp_dep		salary	dept_name d	•	
1	Smith	-1	2018	10	M	3000		10	,
3	Williams	1	2010	10	M	1000	Finance	10	- 1
4	Jones	2	2005	10	F	2000	Finance	10	ĺ
12	Rose	1	2010	20	lm.	4000	Marketing	20	- 1
5	Brown	2	2010	40	j	-1	İIT	40	i
4								_	

+	+	+	+	+	+	+	+	+
emp_id	name	superior_emp_id	year_joined	emp_dept_id {	gender s	alary d	lept_name d	ept_id
+ 1	Smith	+ -1	+ 2018	+ 10	+ М	+ 3000	+ Finance	+ 10
3	Williams	1	2010	10	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2010	20	M	4000	Marketing	20
null	null	null	null	null	null	null	Sales	30
5	Brown	2	2010	40	l	-1	IT	40
6	Brown	2	2010	50	l	-1	null	null
+ emp_id	name	+ superior_emp_id	year_joined	+emp_dept_id {	ender s	alary d	+ ept_name d	+ ept_id
+	+	+	+	+	+	+	+	+
1	Smith		2018	•	lΜ	3000	Finance	10
3	Williams	1	2010	'	M	1000	Finance	10
4	Jones	2	2005	10	F	2000	Finance	10
2	Rose	1	2010	20	M	•	Marketing	20
null	null	null	null	null	null	null	Sales	30

5. PySpark Left Outer Join

Left a.k.a Leftouter join returns all rows from the left dataset regardless of match found on the right dataset when join expression doesn't match, it assigns null for that record and drops records from right where match not found.

From our dataset, "emp_dept_id" 50 doesn't have a record on "dept" dataset hence, this record contains null on "dept" columns (dept_name & dept_id). and "dept_id" 30 from "dept" dataset dropped from the results. Below is the result of the above Join expression.

```
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"left") \
    .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftouter") \
    .show(truncate=False)
```

1	Smith	+ -1	+ 2018	10	+ M	+ 3000	Finance	+ 10	+
2	Rose	1	2010	20	M	4000	Marketing		i
3	Williams	1	2010	10	M	1000	Finance	10	ĺ
4	Jones	2	2005	10	F	2000	Finance	10	
5	Brown	2	2010	40		-1	IT	40	
6	Brown	2	2010	50		-1	null	null	
+	+	+	+	+	+	+	-+	+	+
+ Lomn	id name	leupopion	emp id year joi				-+	•	
Lemb_	_Tu IIallie	Lanber Tor.	emb_ru year_Jor	neu reinp_uep		ii. I satai. y	uept_name u	ept_iu	

1	Smith	-1	2018	10	M	3000	Finance	10	
2	Rose	1	2010	20	M	4000	Marketing	20	
3	Williams	1	2010	10	M	1000	Finance	10	
4	Jones	2	2005	10	F	2000	Finance	10	
5	Brown	2	2010	40		-1	IT	40	
6	Brown	2	2010	50		-1	null	null	

...

6. Right Outer Join

Right a.k.a Rightouter join is opposite of left join, here it returns all rows from the right dataset regardless of math found on the left dataset, when join expression doesn't match, it assigns null for that record and drops records from left where match not found.

From our example, the right dataset "dept_id" 30 doesn't have it on the left dataset "emp" hence, this record contains null on "emp" columns. and "emp_dept_id" 50 dropped as a match not found on left. Below is the result of the above Join expression.

```
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"right") \
    .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"rightouter") \
    .show(truncate=False)
```

emp_ic	-+ d name -+	superior_emp_id	year_joined	emp_dept_id	gender	+ salary 0 +	dept_name d +	+ ept_id +	+
4	Jones	2	2005	10	F	2000	Finance	10	i
3	Williams	1	2010	10	M	1000	Finance	10	l
1	Smith	-1	2018	10	M	3000	Finance	10	
2	Rose	1	2010	20	M	4000	Marketing	20	
null	null	null	null	null	null	null	Sales	30	
5	Brown	2	2010	40	1	-1	IT	40	
+	-+	+	+	+	+	+	+	+	+
+	-+	+	+	. +	+	+	+	+	+
emp_io	d name	superior_emp_id	.,				. – .	. – .	

emp_i	d name	superior_emp_id	dept_name d	dept_id					
+	-+	+	+	+	+	+	+	+	+
4	Jones	2	2005	10	F	2000	Finance	10	
3	Williams	1	2010	10	M	1000	Finance	10	
1	Smith	-1	2018	10	M	3000	Finance	10	
2	Rose	1	2010	20	M	4000	Marketing	20	
null	null	null	null	null	null	null	Sales	30	
5	Brown	2	2010	40		-1	IT	40	
		_							

...

7. Left Semi Join

leftsemi join is similar to inner join difference being leftsemi join returns all columns from the left dataset and ignores all columns from the right dataset. In other words, this join returns columns from the only left dataset for the records match in the right dataset on join expression, records not matched on join expression are ignored from both left and right datasets.

The same result can be achieved using select on the result of the inner join however, using this join would be efficient. Below is the result of the above join expression.

empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftsemi") \
 .show(truncate=False)

+	+	+	+	+	+	+	-+
emp	_id name +	superior	·_emp_id year_joir	ned emp_dept	t_id gende	r salary +	-+
1	Smith	-1	2018	10	M	3000	İ
3	William	ns 1	2010	10	M	1000	
4	Jones	2	2005	10	F	2000	
2	Rose	1	2010	20	M	4000	
5	Brown	2	2010	40		-1	

1.

```
8. Left Anti Join
leftanti join does the exact opposite of the leftsemi, leftanti join returns only columns from the left dataset for non-matched records.

Yields below output
...
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftanti") \
..show(truncate=False)
```

```
+----+
|emp_id|name |superior_emp_id|year_joined|emp_dept_id|gender|salary|
+----+
|6 |Brown|2 |2010 |50 | |-1 |
```

```
Joins are not complete without a self join, Though there is no self-join type available, we can use any of the above-explained join types to join DataFrame to itself. below example use inner self join.

Here, we are joining emp dataset with itself to find out superior emp_id and name for all employees.

'''

from pyspark.sql.functions import col

empDF.alias("emp1").join(empDF.alias("emp2"), \

col("emp1.superior_emp_id") == col("emp2.emp_id"), "inner") \

.select(col("emp1.emp_id"),col("emp1.name"), \

col("emp2.emp_id").alias("superior_emp_id"), \

col("emp2.name").alias("superior_emp_name")) \

.show(truncate=False)
```

emp_id name	+ superior_emp_id +	++ superior_emp_name
1- 1		Smith
3 Williams 4 Jones	•	Smith Rose
5 Brown 6 Brown	•	Rose Rose
+	· +	· +

```
4. Using SQL Expression
Since PySpark SQL support native SQL syntax, we can also write join operations after creating temporary tables on DataFrames and use these tables on spark.sql().

...

empDF.createOrReplaceTempView("EMP")

deptDF.createOrReplaceTempView("DEPT")

joinDF = spark.sql("select * from EMP e, DEPT d where e.emp_dept_id == d.dept_id") \
    .show(truncate=False)

joinDF2 = spark.sql("select * from EMP e INNER JOIN DEPT d ON e.emp_dept_id == d.dept_id") \
    .show(truncate=False)
```

-+d name	+ superior_emp_id	+ year_joined	+ emp_dept_id	+ gender :	+ salary d	+ lept_name d	+ ept_id	÷
-+ Smith	+	+ 2018	+ 10	+ M	+ 3000	+ Finance	+ 10	ŀ
Williams	1	2010	10	Im	1000	Finance	10	i
Jones	2	2005	10	ļF	2000	Finance	10	i
Rose	1	2010	20	M	4000	Marketing	20	İ
Brown	2	2010	40	İ	-1	IT	40	İ
-+d name	+ superior_emp_id	+ year_joined	+ emp_dept_id	+ gender :	+salary c	+ lept_name d	+ ept_id	۲
-+ Smi+h	+ _1	+ 2018	+ 10	+ Ім	+ 13000	t Finance	+ 10	ř L
	!	!	!	•			!	ï
:		:	:	•	1		:	i
	•	•	•					i
Brown	2	2010	40	i	-1	IT	40	İ
	Williams Jones Rose Brown	Smith -1 Williams 1 Jones 2 Rose 1 Brown 2	Smith -1	Smith -1	Smith -1	Smith -1	Smith -1	Smith -1

```
6. PySpark SQL Join Complete Example
import pyspark
from pyspark.sql import SparkSession
from pyspark.sql.functions import col
spark = SparkSession.builder.appName('SparkByExamples.com').getOrCreate()
emp = [(1,"Smith",-1,"2018","10","M",3000), \
    (2,"Rose",1,"2010","20","M",4000), \
    (3,"Williams",1,"2010","10","M",1000), \
    (4,"Jones",2,"2005","10","F",2000), \
    (5,"Brown",2,"2010","40","",-1), \
     (6, "Brown", 2, "2010", "50", "", -1) \
empColumns = ["emp_id","name","superior_emp_id","year_joined", \
       "emp_dept_id","gender","salary"]
empDF = spark.createDataFrame(data=emp, schema = empColumns)
empDF.printSchema()
empDF.show(truncate=False)
dept = [("Finance",10), \
    ("Marketing",20), \
    ("Sales",30), \
    ("IT",40) \
 ]
deptColumns = ["dept_name","dept_id"]
deptDF = spark.createDataFrame(data=dept, schema = deptColumns)
deptDF.printSchema()
deptDF.show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"inner") \
     .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"outer") \
    .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"full") \
    .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"fullouter") \
    .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"left") \
    .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftouter") \
  .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"right") \
   .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"rightouter") \
   .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftsemi") \
   .show(truncate=False)
empDF.join(deptDF,empDF.emp_dept_id == deptDF.dept_id,"leftanti") \
   .show(truncate=False)
empDF.alias("emp1").join(empDF.alias("emp2"), \
    col("emp1.superior_emp_id") == col("emp2.emp_id"),"inner") \
    .select(col("emp1.emp_id"),col("emp1.name"), \
     col("emp2.emp_id").alias("superior_emp_id"), \
     col("emp2.name").alias("superior_emp_name")) \
  .show(truncate=False)
```

emp_	_id name	superior_	emp_id year_joi	ned emp_dept	_id gende	r salary
1	Smith	-1	2018	10	M	3000
2	Rose	1	2010	20	M	4000
3	Williams	1	2010	10	M	1000
4	Jones	2	2005	10	F	2000
5	Brown	2	2010	40		-1
6	Brown	2	2010	50		-1

root

```
PySpark distinct() function is used to drop/remove the duplicate rows (all columns) from DataFrame and dropDuplicates() is used
to drop rows based on selected (one or multiple) columns. In this article, you will learn how to use distinct() and
dropDuplicates() functions with PySpark example.
Before we start, first let's create a DataFrame with some duplicate rows and values on a few columns. We use this DataFrame to
demonstrate how to get distinct multiple columns.
On the above table, record with employer name James has duplicate rows. As you notice we have 2 rows that have duplicate values
on all columns and we have 4 rows that have duplicate values on department and salary columns.
import pyspark
from pyspark.sql import SparkSession
from pyspark.sql.functions import expr
spark = SparkSession.builder.appName('SparkByExamples.com').getOrCreate()
data = [("James", "Sales", 3000), \
    ("Michael", "Sales", 4600), \
    ("Robert", "Sales", 4100), \
    ("Maria", "Finance", 3000), \
    ("James", "Sales", 3000), \
   ("Scott", "Finance", 3300), \
    ("Jen", "Finance", 3900), \
    ("Jeff", "Marketing", 3000), \
    ("Kumar", "Marketing", 2000), \
    ("Saif", "Sales", 4100) \
columns= ["employee_name", "department", "salary"]
df = spark.createDataFrame(data = data, schema = columns)
df.printSchema()
df.show(truncate=False)

    Get Distinct Rows (By Comparing All Columns)

On the above DataFrame, we have a total of 10 rows with 2 rows having all values duplicated, performing distinct on this
DataFrame should get us 9 after removing 1 duplicate row.
distinct() function on DataFrame returns a new DataFrame after removing the duplicate records. This example yields the below
output.
#Distinct
distinctDF = df.distinct()
print("Distinct count: "+str(distinctDF.count()))
distinctDF.show(truncate=False)
Alternatively, you can also run dropDuplicates() function which returns a new DataFrame after removing duplicate rows.
#Drop duplicates
df2 = df.dropDuplicates()
print("Distinct count: "+str(df2.count()))
df2.show(truncate=False)
2. PySpark Distinct of Selected Multiple Columns
PySpark doesn't have a distinct method which takes columns that should run distinct on (drop duplicate rows on selected multiple
columns) however, it provides another signature of dropDuplicates() function which takes multiple columns to eliminate
duplicates.
Note that calling dropDuplicates() on DataFrame returns a new DataFrame with duplicate rows removed.
Yields below output. If you notice the output, It dropped 2 records that are duplicate.
#Drop duplicates on selected columns
dropDisDF = df.dropDuplicates(["department","salary"])
print("Distinct count of department salary : "+str(dropDisDF.count()))
dropDisDF.show(truncate=False)
```

Distinct count: 9

In this PySpark SQL article, you have learned distinct() method which is used to get the distinct values of rows (all columns) and also learned how to use dropDuplicates() to get the distinct and finally learned using dropDuplicates() function to get distinct of multiple columns.

root |-- employee_name: string (nullable = true) |-- department: string (nullable = true) |-- salary: long (nullable = true) |employee_name|department|salary| |Sales |3000 | James |4600 Sales Michael Robert |Sales 4100 |Finance |3000 | lMaria James |Sales | 3000 | |Finance |3300 | Scott |Jen |Finance |3900 |Jeff |Marketing |3000 |Kumar |Marketing |2000 | Saif |Sales | 4100 | +----+

```
PySpark JSON Functions with Examples
PySpark JSON functions are used to query or extract the elements from JSON string of DataFrame column by path, convert it to
struct, mapt type e.t.c, In this article, I will explain the most used JSON SQL functions with Python examples.
1. PySpark JSON Functions
from_json() - Converts JSON string into Struct type or Map type.
to_json() - Converts MapType or Struct type to JSON string.
json_tuple() - Extract the Data from JSON and create them as a new columns.
get_json_object() - Extracts JSON element from a JSON string based on json path specified.
schema of json() - Create schema string from JSON string
1.1. Create DataFrame with Column contains JSON String
In order to explain these JSON functions first, let's create DataFrame with a column contains JSON string.
from pyspark.sql import SparkSession,Row
spark = SparkSession.builder.appName('SparkByExamples.com').getOrCreate()
jsonString="""{"Zipcode":704,"ZipCodeType":"STANDARD","City":"PARC PARQUE","State":"PR"}"""
df=spark.createDataFrame([(1, jsonString)],["id","value"])
df.show(truncate=False)
2.1. from_json()
PySpark from_json() function is used to convert JSON string into Struct type or Map type. The below example converts JSON string
to Map key-value pair. I will leave it to you to convert to struct type. Refer, Convert JSON string to Struct type column.
#Convert JSON string column to Map type
from pyspark.sql.types import MapType,StringType
from pyspark.sql.functions import from_json
df2=df.withColumn("value",from_json(df.value,MapType(StringType()),StringType())))
df2.printSchema()
df2.show(truncate=False)
2.2. to json()
to json() function is used to convert DataFrame columns MapType or Struct type to JSON string. Here, I am using df2 that created
from above from_json() example.
from pyspark.sql.functions import to_json,col
df2.withColumn("value",to_json(col("value"))) \
   .show(truncate=False)
2.3. ison tuple()
Function json_tuple() is used the query or extract the elements from JSON column and create the result as a new columns.
from pyspark.sql.functions import json_tuple
df.select(col("id"),json_tuple(col("value"),"Zipcode","ZipCodeType","City")) \
    .toDF("id","Zipcode","ZipCodeType","City") \
    .show(truncate=False)
2.4. get_json_object()
get_json_object() is used to extract the JSON string based on path from the JSON column.
from pyspark.sql.functions import get_json_object
df.select(col("id"),get_json_object(col("value"),"$.ZipCodeType").alias("ZipCodeType")) \
    .show(truncate=False)
```

```
2.5. schema_of_json()
 Use schema_of_json() to create schema string from JSON string column.
 from pyspark.sql.functions import schema_of_json,lit
 schemaStr=spark.range(1) \
     .select(schema_of_json(lit("""{"Zipcode":704,"ZipCodeType":"STANDARD","City":"PARC PARQUE","State":"PR"}"""))) \
     .collect()[0][0]
 print(schemaStr)
lid |value
|1 |{"Zipcode":704,"ZipCodeType":"STANDARD","City":"PARC PARQUE","State":"PR"}|
root
|-- id: long (nullable = true)
|-- value: map (nullable = true)
| |-- key: string
    |-- value: string (valueContainsNull = true)
|id |value
|1 |{Zipcode -> 704, ZipCodeType -> STANDARD, City -> PARC PARQUE, State -> PR}|
|id |value
|-- employee_name: string (nullable = true)
|-- department: string (nullable = true)
|-- salary: long (nullable = true)
|employee_name|department|salary|
+-----
James
           |Sales |3000 |
Michael
           Sales
                    4600
         |Sales | 4100 |
Robert
Maria
          |Finance |3000 |
                     |3000 |
          |Sales
lJames
Scott
           |Finance |3300 |
|Jen
            |Finance
                     3900
|Jeff
           |Marketing |3000 |
Kumar
          |Marketing |2000 |
Saif
          |Sales |4100 |
+----+
+----+
```

databricks Pyspark join types and dataframes Notebook 2023-

(https://databricks.com)

09-04 11:03:40

```
root
|-- emp_id: long (nullable = true)
|-- name: string (nullable = true)
|-- superior_emp_id: long (nullable = true)
|-- year_joined: string (nullable = true)
|-- emp_dept_id: string (nullable = true)
|-- gender: string (nullable = true)
|-- salary: long (nullable = true)
+----+
|emp_id|name |superior_emp_id|year_joined|emp_dept_id|gender|salary|
     |Smith|-1
                       2018
                                 10
                                                3000
12
     |Rose |1
                       2010
                                 120
                                          M
                                                4000
     Jones 2
                       2005
                                 10
                                          |F
                                                2000
15
     |Brown|2
                       2016
                                 40
                                                |-1
                       2010
                                 |50
                                          |Brown|2
                                                |-1
+----+
```

dept_nam	e dept_	id
+	-+	+
Finance	10	
Marketin	g 20	- 1
Sales	30	- 1
IT	40	- 1
+	-+	+

	_id name superior	_emp_id year_jo	•	•		•		+ d
+ 1	+ Smith -1	2018	10	+ M	-++ 3000	+- Finance	10	+
4	Jones 2	2005	10	F	2000	Finance	10	
2	Rose 1	2010	20	M	4000	Marketing	20	
5	Brown 2	2016	40		-1	IT	40	
+	+	+	+	+	-++	+-		+

id 	-+			_	L		
 + 1	Smith -1	2018	10	M	3000	Finance	10
 4	Jones 2	2005	10	F	2000	Finance	10
 2 	Rose 1	2010	20	M	4000	Marketing	g 20
null	null null	null	null	null	null	Sales	30
5	Brown 2	2016	40	1	-1	IT	40
6	Brown 2	2010	50	I	-1	null	null

+ emp_id	++ name superior_emp_io ++	+ d year_joined	+ emp_dept_id	+- gender +	salary	dept_name c	
1	Smith -1	2018	10	M	3000	Finance	10
2	Rose 1	2010	20	M	4000	Marketing	20
4	Jones 2	2005	10	F	2000	Finance	10
5	Brown 2	2016	40		-1	IT	40
6	Brown 2	2010	50		-1	null	null
+ + emp_id	++ ++ name superior_emp_io	d year_joined	•	+- +- gender	•	dept_name c	
1 2 4	Smith -1 Rose 1 Jones 2	2018 2010 2005	10 20 10	M M M F	3000 4000 2000	 Marketing	10 20 10

5	Brown 2	2016	40	-1	IT	40	
6	Brown 2	2010	50	-1	null	null	
				 L	_	т.	_

	1- 1-	1	1			1 - 4	
4	Jones 2	2005	10	F	2000	Finance	10
1	Smith -1	2018	10	M	3000	Finance	10
2	Rose 1	2010	20	M	4000	Marketing	20
null	null null	null	null	null	null	Sales	30
5	Brown 2	2016	40		-1	IT	40
 emp_i	-++d name superior_	+ + emp_id year_joi	+ + ined emp_dept	+ +:_id gender	+ + salary	+- +- dept_name 0	 dept_
+	-+	+	+	+	+	·+-	
 4	-++ Jones 2	2005	10	+ F	2000	Finance	10
 4 1	-++ Jones 2 Smith -1	2005 2018	10 10	F M	2000 3000	Finance Finance	10 10
 4	-++ Jones 2 Smith -1 Rose 1	2005	10	+ F	2000	Finance	10 10
 4 1	-++ Jones 2 Smith -1	2005 2018	10 10	F M	2000 3000	Finance Finance	10 10

|emp_id|name |superior_emp_id|year_joined|emp_dept_id|gender|salary|

1	Smith -1	2018	10	M	3000	
4	Jones 2	2005	10	F	2000	
2	Rose 1	2010	20	M	4000	
5	Brown 2	2016	40	1	-1	
+	+		+	+	+	. +

emp_id	++	id year_joined	d emp_dept_id	gender	salary	
	Brown 2	2010	50		-1 -1	

emp_id	name superior_emp_id	superior.emp_name
2 4	Rose 1 Jones 2	Smith Rose
5 6	Brown 2 Brown 2	Rose
+	++	+

+i	id name superior	_emp_id year_joi	+ ined emp_dep	ot_id gende	-++ er salary -+	dept_name	dept_i	+ .d +
1	Smith -1	2018	10	M	3000	Finance	10	
4	Jones 2	2005	10	F	2000	Finance	10	I
2	Rose 1	2010	20	M	4000	Marketing	g 20	
5	Brown 2	2016	40		-1	IT	40	
+i emp_i	id name superior	_emp_id year_joi	+ined emp_dep	•		·	dept_i	+ .d _
1	Smith -1	2018	10	M	3000	Finance	10	,
4	Jones 2	2005	10	F	2000	Finance	10	i
2	Rose 1	2010	20	M	4000	Marketing	g 20	
5	Brown 2	2016	40		-1	IT	40	

root |-- employee_name: string (nullable = true) |-- department: string (nullable = true)

|-- salary: long (nullable = true)

++					
employee_name department salary					
+	+	+	+		
James	Sales	3000			
Michael	Sales	4600			
Robert	Sales	4100			
Maria	Finance	3000			
James	Sales	3000			
Scott	Finance	3300			
Jen	Finance	3988			
Jeff	Marketing	3000			
Kumar	Marketing	2000			
Saif	Sales	4100			

+----+

Disnict count	: 9					
+	+	++				
employee_name department salary						
+	+	++				
James	Sales	3000				
Michael	Sales	4600				
Robert	Sales	4100				
Maria	Finance	3000				
Scott	Finance	3300				
Jen	Finance	3988				
Jeff	Marketing	3000				
Kumar	Marketing	2000				
Saif	Sales	4100				
+	+	++				

Disnict count	: 9		
+	-+	+	+
employee_name department salary			
++			
James	Sales	3000	
Michael	Sales	4600	
Robert	Sales	4100	
Maria	Finance	3000	
Scott	Finance	3300	
Jen	Finance	3988	
Jeff	Marketing	3000	
Kumar	Marketing	2000	
Saif	Sales	4100	
++			

Disnict count	Selected	items: 8		
++				
employee_name department salary				
++				
Maria	Finance	3000		
Scott	Finance	3300		
Jen	Finance	3988		
Kumar	Marketing	g 2000		
Jeff	Marketing	g 3000		
James	Sales	3000		
Robert	Sales	4100		
Michael	Sales	4600		
+	-+	-++		

♦ databricksSpark SQL G1 and G2 Notebook 2023-08-28 12:31:23

```
(https://databricks.com)
   val sqlConstext = new org.apache.spark.sql.SQLContext(sc)
 command-4344394284453045:1: warning: constructor SQLContext in class SQLContext is deprecated (since 2.0.0): Use SparkSession.builde
 val sqlConstext = new org.apache.spark.sql.SQLContext(sc)
 sqlConstext: org.apache.spark.sql.SQLContext = org.apache.spark.sql.SQLContext@508225e8
   val a = sc.parallelize(1 to 10)
 a: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[17] at parallelize at command-4344394284453046:1
   val b = a.map(x=>(x , x+1))
 b: org.apache.spark.rdd.RDD[(Int, Int)] = MapPartitionsRDD[18] at map at command-4344394284453047:1
   h.collect
 res0: Array[(Int, Int)] = Array((1,2), (2,3), (3,4), (4,5), (5,6), (6,7), (7,8), (8,9), (9,10), (10,11))
   val df = b.toDF("First", "Second")
 df: org.apache.spark.sql.DataFrame = [First: int, Second: int]
    df.show
 |First|Second|
             2
      1|
      2|
             3|
             4
      3|
             5|
      51
             6
      6
             7
      71
             8
      8|
             9|
      9|
            10|
     10
            11
    val a = List(("Tom", 5),("Jerry", 2),("Donald", 7))
 a: List[(String, Int)] = List((Tom,5), (Jerry,2), (Donald,7))
    val df = a.toDF("Name", "Age")
 df: org.apache.spark.sql.DataFrame = [Name: string, Age: int]
    df.show
 | Name|Age|
```

```
| Tom| 5|
| Jerry| 2|
|Donald| 7|
  val a = Seq(("Tom", 5),("Jerry", 2),("Donald", 7))
a: Seq[(String, Int)] = List((Tom,5), (Jerry,2), (Donald,7))
  val df = a.toDF("Name", "Age")
df: org.apache.spark.sql.DataFrame = [Name: string, Age: int]
 df.show
| Name|Age|
+----+
| Tom| 5|
| Jerry| 2|
|Donald| 7|
  df.registerTempTable("Cartoon")
command-4344394284453057:1: warning: method registerTempTable in class Dataset is deprecated (since 2.0.0): Use createOrReplaceTempV
iew(viewName) instead.
df.registerTempTable("Cartoon")
  df.createOrReplaceTempView("Cartoon")
  sqlContext.sql("select * from Cartoon where Name = 'Tom'").show
+----+
|Name|Age|
+---+
| Tom| 5|
+---+
  sqlContext.sql("select * from Cartoon").show
+----+
| Name | Age |
+----+
| Tom| 5|
| Jerry| 2|
|Donald| 7|
+----+
```

```
sqlContext.sql("select count(*) from Cartoon").show
|count(1)|
+----+
| 3|
+----+
  // questions : to create a json file, upoad it open dbfs and perform the following operations on it.
  // printSchema()
  // select the query with all the names
  // filter and identify age > 23
  // groupBy Age Count it and show it
  // how ro read file
  // var df1 = spark.read.format("json").load("dbfs:/FileStore/shared_uploads/....../....json")
  var df1 = spark.read.format("json").load("/FileStore/tables/file.json")
df1: org.apache.spark.sql.DataFrame = [_corrupt_record: string]
  display(df1)
  AnalysisException: Since Spark 2.3, the queries from raw JSON/CSV files are disallowed when the
referenced columns only include the internal corrupt record column
(named _corrupt_record by default). For example:
\verb|spark.read.schema|.csv(file).filter($"\_corrupt\_record".isNotNull).count()|\\
and spark.read.schema(schema).csv(file).select("_corrupt_record").show().
Instead, you can cache or save the parsed results and then send the same query.
For example, val df = spark.read.schema(schema).csv(file).cache() and then
df.filter($"_corrupt_record".isNotNull).count().
  val df1 = spark.read.format("json").load("dbfs:/FileStore/shared_uploads/devjethva234@gmail.com/emp_1.json")
df1: org.apache.spark.sql.DataFrame = [age: string, id: string ... 1 more field]
  df1.show
|age| id|
            name
| 25|1201|
               om
25 | 1202 |
               some
| 25|1203|
              thing|
| 25|1204|different|
25 | 1205 |
             going
| 25|1206|
                on
25 | 1207 |
             kavan
+---+---+-----+
  AnalysisException: [TABLE_OR_VIEW_NOT_FOUND] The table or view `df1` cannot be found. Verify the spelling and correctness of the s
```

chema and catalog.

If you did not qualify the name with a schema, verify the current_schema() output, or qualify the name with the correct schema and ca

talog.
To tolerate the error on drop use DROP VIEW IF EXISTS or DROP TABLE IF EXISTS.; line 1 pos 17;
'Project ['name]
+- 'UnresolvedRelation [df1], [], false

Hadoop Installation

Prepared by: VASAVA VIPULKUMAR

DINESHBHAI (23MDS003)

Branch: MTech (Data Science)

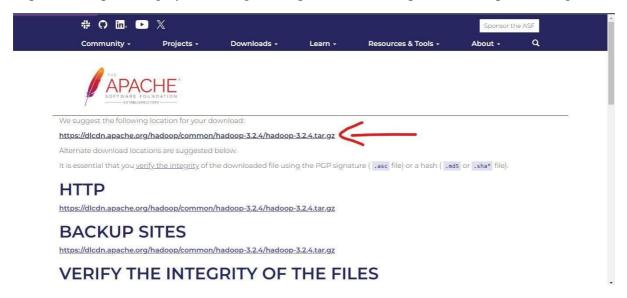
Prerequisites

- 1. Java 8 runtime environment (JRE): Hadoop 3 requires a Java 8 installation.
- 2. Java 8 development Kit (JDK)
- 3. To unzip downloaded Hadoop binaries, we should install 7zip.

❖ Download Hadoop binaries

The first step is to download Hadoop binaries from the official website. we need to install Hadoop 3.2.4 The binary package size is about 470 MB.

https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-3.2.4/hadoop-3.2.4.tar.gz

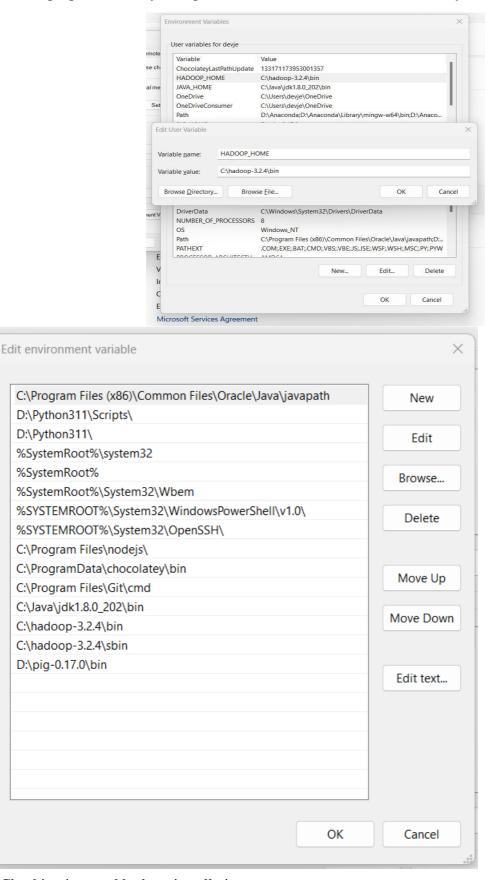


After finishing the file download, we should unpack the package using 7zip. First, we should extract the hadoop-3.2.1.tar.gz library, and then, we should unpack the extracted tar file.

Setting up environment variables

After installing Hadoop and its prerequisites, we should configure the environment variables to define Hadoop and Java default paths.

To edit environment variables, go to Control Panel > System and Security > System (or right-click > properties on My Computer icon) and click on the "Advanced system settings" link.



Checking java and hadoop installation:

Microsoft Windows [Version 10.0.22621.2715]
(c) Microsoft Corporation. All rights reserved.

C:\Users\devje>hadoop version
Hadoop 3.2.4
Source code repository Unknown -r 7e5d9983b388e372fe640f21f048f2f2ae6e9eba
Compiled by ubuntu on 2022-07-12T11:58Z
Compiled with protoc 2.5.0
From source with checksum ee031c16fe785bbb35252c749418712
This command was run using /C:/hadoop-3.2.4/share/hadoop/common/hadoop-common-3.2.4.jar

C:\Users\devje>

***** Configuring Hadoop cluster

There are four files we should alter to configure Hadoop cluster:

1. %HADOOP_HOME%\etc\hadoop\hdfs-site.xml

As we know, Hadoop is built using a master-slave paradigm. Before altering the HDFS configuration file, we should create a directory to store all master node (name node) data and another one to store data (data node). In this example, we created the following directories:

D:\hadoop\hadoop-3.2.4\data\namenode

D:\hadoop\hadoop-3.2.4\data\datanode

Now, let's open "hdfs-site.xml" file located in "%HADOOP_HOME%\etc\hadoop" directory, and we should add the following properties within the <configuration></configuration> element:

```
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value> file:/// D:\hadoop\hadoop-3.2.4\data\namenode</value>
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value> file:/// D:\hadoop\hadoop-3.2.4\data\datanode</value>
</property>
<property>
<property></property></property></property>
```

Note that we have set the replication factor to 1 since we are creating a single node cluster.

2. %HADOOP_HOME%\etc\hadoop\core-site.xml

Now, we should configure the name node URL adding the following XML code into the <configuration></configuration> element within "core-site.xml":

```
<name>fs.default.name
```

```
<value>hdfs://localhost:9820</value>
```

3. %HADOOP_HOME%\etc\hadoop\mapred-site.xml

Now, we should add the following XML code into the <configuration></configuration> element within "mapred-site.xml"

```
<name>mapreduce.framework.name
```

4. %HADOOP_HOME%\etc\hadoop\yarn-site.xml

Now, we should add the following XML code into the <configuration></configuration> element within "yarn-site.xml"

```
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
<description>Yarn Node Manager Aux Service</description>
```

***** Formatting Name node

After finishing the configuration, let's try to format the name node using the following command:

hdfs namenode -format

Starting Hadoop services

Now, we will open PowerShell, and navigate to "%HADOOP_HOME%\sbin" directory or just open cmd as admin. Then we will run the following command to start the Hadoop nodes:

start-all

This will run both dfs and yarn, must have to run all 4 terminal, no one have to shutdown, than installation was successful also check this with 'jps' it display all running services.

```
Administrator. Command Prompt

Microsoft Windows [Version 10.0.19045.3570]

(c) Microsoft Corporation. All rights reserved.

C:\WINDOWS\system32>start-all
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\WINDOWS\system32>jps
8204 NameNode
3416 DataNode
4408 ResourceManager

6072 NodeManager

C:\WINDOWS\system32>

\[
\tilde{\text{MINDOWS\system32}}
\]

\[
\tilde{\text{NINDOWS\system32}}
\]

\[
\tilde{\text{MINDOWS\system32}}
\]

\[
\tilde{\text{NINDOWS\system32}}
\]

\[
\tilde{\text{NINDOWS\system32}}
\]

\[
\tilde{\text{MINDOWS\system32}}
\]

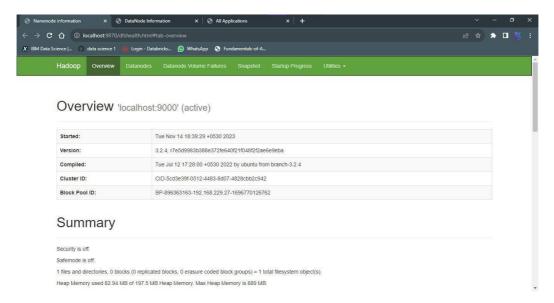
\[
\t
```

```
23-11-14 16:56:32,045 INFO impl.FsDatasetImpl: Total time to add all replicas to map for block pool BP-896363163-192.1
229.27-1696770126762: 10ms
23-11-14 16:56:32,047 INFO checker.ThrottledAsyncChecker: Scheduling a check for D:\hadoop\hadoop-3.2.4\data\datanode
23-11-14 16:56:32,066 INFO checker.DatasetVolumeChecker: Scheduled health check for volume D:\hadoop\hadoop-3.2.4\data
      -11-14 16:56:32,366 INFO checker.DatasetVolumeChecker: Scheduled health check for volume D:\hadoop\hadoop-3.2.4\data anode anode anode anode anode anode anode anode anode anode anode anode anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote anote ano
          11-14 17:45:52,768 WARN util.JvmPauseMonitor: Detected pause in JVM or host machine (eg GC): pause of approximately
167ms
3 detected
11-14 18:14:29,932 WARN util.JvmPauseMonitor: Detected pause in JVM or host machine (eg GC): pause of approximately
     s.
5 detected
11-14 18:18:46,755 WARN util.JvmPauseMonitor: Detected pause in JVM or host machine (eg GC): pause of approximately
                 11039091
14 16:56:33.711 INFO security.NMTokenSecretManagerInNM: Rolling master-key for container-tokens, got key with id
                 24
14 16:56:33,714 INFO nodemanager.NodeStatusUpdaterImpl: Registered with ResourceManager as DESKTOP-0007548:55891
tal resource of Kmemory:8192, vCores:8>
14 17:86:24,389 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public
: 0, Private Deleted: 0
14 17:45:52,842 WARN util.JvmPauseMonitor: Detected pause in JVM or host machine (eg GC): pause of approximately
             Sms
detected
-14 17:45:53,425 INFO retry.RetryInvocationHandler: java.io.IOException: DestHost:destPort 0.0.0.0:8031 , LocalHo
-14 17:45:53,425 INFO retry.RetryInvocationHandler: java.io.IOException: DestHost:destPort 0.0.0.0:8031 , LocalHo
-10 losed by the remote host, while invoking ResourceTrackerPBC.InentImpl.nodeHeartDeat over null. Retrying after slee
- 30000Mns. Current retry count: 0.
-14 17:54:22,408 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public
d: 0, Private Deleted: 0
-14 18:14:29,932 WARN util.JvmPauseMonitor: Detected pause in JVM or host machine (eg GC): pause of approximately
ms
              -14 18:18:46,755 WARN util.JymPauseMonitor: Detected pause in JVM or host machine (eg GC): pause of approximately
            ms
detected
-14 18:19:22,362 INFO localizer.ResourceLocalizationService: Cache Size Before Clean: 0, Total Deleted: 0, Public
di 0, Private Deleted: 0
```

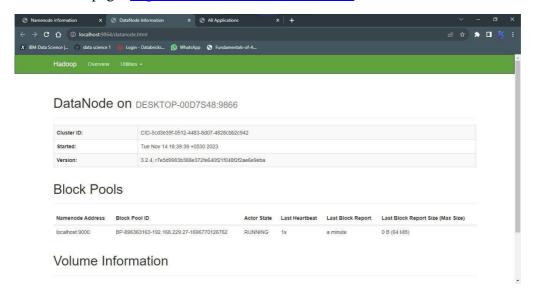
***** Hadoop Web UI

There are three web user interfaces to be used:

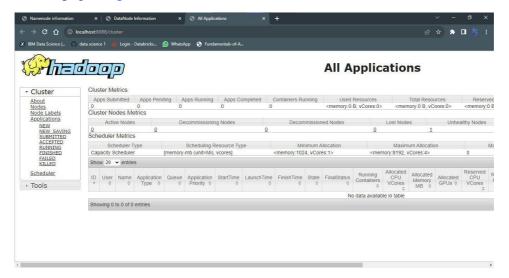
Name node web page: http://localhost:9870/dfshealth.html



Data node web page: http://localhost:9864/datanode.html



Yarn web page: http://localhost:8088/cluster



Pig Installation

Prepared by:

VASAVA VIPULKUMAR DINESHBHAI (23MDS003)

Branch: MTech (Data Science)

Prerequisites

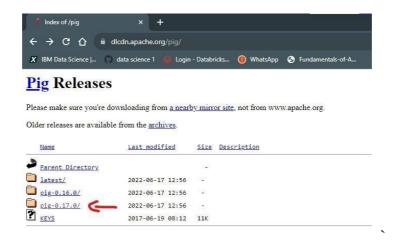
- 1. Java 8 development Kit (JDK)
- 2. Hadoop 3.2.4 must be installed.
- 3. 7zip

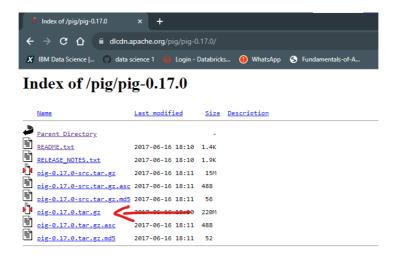
Download PIG file

The first step is to download pig form the apache official server, it can be downloaded at following link:

https://dlcdn.apache.org/pig/

under this link go to the pig-0.17.0/ and then download <u>pig-0.17.0.tar.gz</u> this file. The package size is about 220MB.

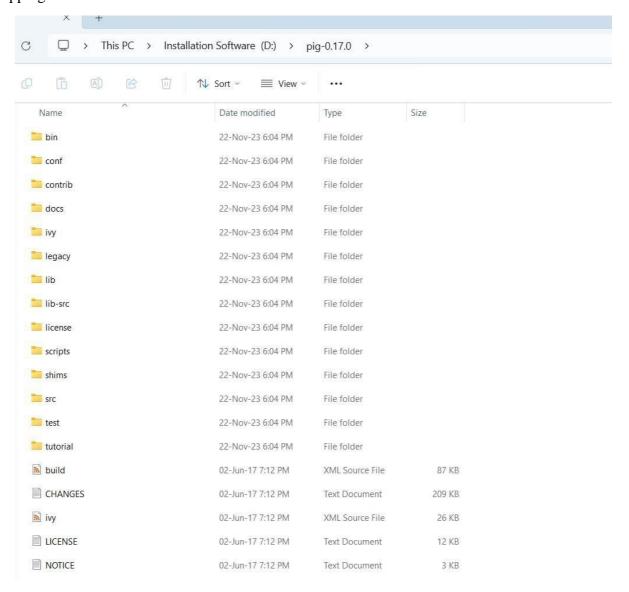




Unzip and Install PIG

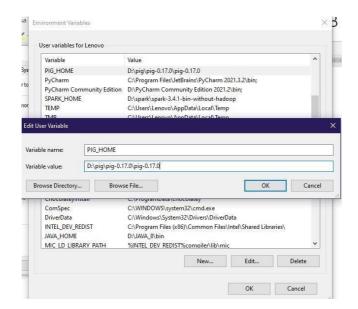
After Downloading the PIG, we need to Unzip the pig-0.17.0.tar.gz file.

After unzipping the folder it will be look like this:

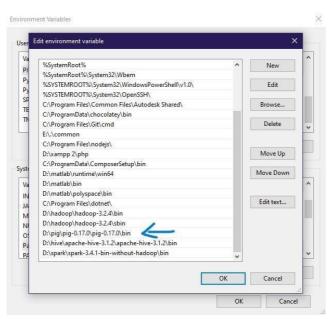


Setting Up Environment Variables

- Another important step in setting up a work environment is to set your Systems environment variable.
- Go to Control Panel > System > click on the "Advanced system settings" link to edit environment variables.
- Open environment Variable and click on "New" in "User Variable".
- On clicking "New", we get the below screen.



• The last step in setting the Environment variable is setting Path in System Variable.



Verify the Paths

Now we need to verify that what we have done is correct and reflecting.

- Open a NEW Command Window
- Run following commands

echo %PIG_HOME%

```
Select Command Prompt

Microsoft Windows [Version 10.0.19045.3570]
(C) Microsoft Corporation. All rights reserved.

C:\Users\Lenovo>echo %PIG_HOME%
D:\pig\pig-0.17.0\pig-0.17.0

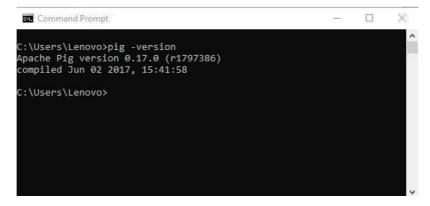
C:\Users\Lenovo>
```

❖ Verifying Setup

- We are done with setting up the PIG on our System.
- Now we need to check if everything works smoothly...
- Open a cmd window, run the below command to test the connection and PIG.

pig -version

Upon running the command we should get the version of PIG. i.e 0.17.0 in our case.



Don't worry some of us can get the below error after running pig -version

'-Xmx1000M' is not recognized as an internal or external command, operable program or batch file.

To resolve this we will need to perform the following steps:-

- 1. Open the pig. cmd file in edit mode. We can find the file in the bin folder.
- 2. Now we need to change the value of the HADOOP_BIN_PATH

3. Save the file.

The next step is to verify the setup once again. So, we need to execute the pigversion command once again.

Starting PIG

Now we need to start a new Command Prompt remember to run it as administrator to avoid permission issues and execute the below commands:

pig

```
C:\Users\Lenovo>pig -version
Apache Pig version 0.17.0 (r1797386)
compiled Jun 02 2017, 15:41:58

C:\Users\Lenovo>pig
2023-11-15 15:41:19,101 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
2023-11-15 15:41:19,156 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE
2023-11-15 15:41:19,157 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType
2023-11-15 15:41:20,362 [main] INFO org.apache.pig.Main - Apache Pig version 0.17.0 (r179738
6) compiled Jun 02 2017, 15:41:58
2023-11-15 15:41:20,362 [main] INFO org.apache.pig.Main - Logging error messages to: D:\hado
op\hadoop-3.2.4\logs\pig_1700043080280.log
2023-11-15 15:41:20,513 [main] INFO org.apache.pig.impl.util.Utils - Default bootup file C:\
Users\Lenovo.pigbootup not found
2023-11-15 15:41:21,724 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapre
d.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2023-11-15 15:41:21,724 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecution
Engine - Connecting to hadoop file system at: hdfs://localhost:9000
2023-11-15 15:41:24,327 [main] INFO org.apache.pig.PigServer - Pig Script ID for the session
: PIG-default-14a1c850-bc4e-410d-aa72-61577af03dd4
2023-11-15 15:41:24,328 [main] WARN org.apache.pig.PigServer - ATS is disabled since yarn.ti
meline-service.enabled set to false
grunt>
```

We can see grunt> once the pig starts.

Kafka Installation

Prepared by:

VASAVA VIPULKUMAR

DINESHBHA (23MDS003)

Branch: MTech. (Data Science)

***** Prerequisites

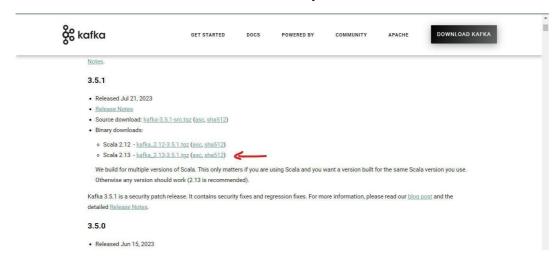
- 1. Java 8 development Kit (JDK)
- 2. 7zip

* Download Kafka file

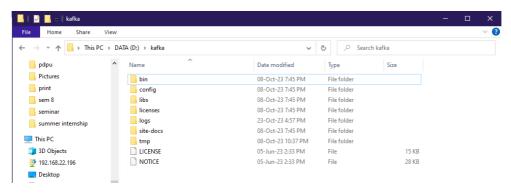
The first step is to download pig form the apache official server, it can be downloaded at following link:

https://kafka.apache.org/downloads

In this link download scala 2.13 version under binary downloads.



After downloading unzip this folder so it look like this



***** Configure files

Now go to kafka>config and open server.properties file, we have to change log directory path to save log on specific location.

As shown in above image set log.dirs path to your specific location.

Now for the same open zookeeper.properties and edit dataDir path to specific location as shown below image.

```
File Edit Selection View Go Run ...  

D: Search

D: Sarka: Sconfig > 2 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties ×

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3 zookeeper.properties x

D: Skarka: Sconfig > 3
```

* Run kafka

To run kafka server first have to start zookeeper server. Now locate to kafka folder where you have bin folder and then open cmd at this place and run following command:

.\bin\windows\zookeeper-server-start.bat .\config\zookeeper.properties

If you seen this zookeeper than server started successfully.

Now we have to start kafka server, so open new terminal at same place and run following command:

.\bin\windows\kafka-server-start.bat .\config\server.properties

Kafka works on producer and consumer model so first we have to create one topic, open new terminal at same place and create a topic by following command:

.\bin\windows\kafka-topics.bat --create --topic {topic name} --bootstrap-server localhost:9092

Write topic name of your choice don't take space or any special character in topic name it consider only one word of name.

After successfully creating a topic, lets start producer by this command, give same topic name that created:

.\bin\windows\kafka-console-producer.bat --topic {topic name} --bootstrap-server localhost:9092

For the consumer open new terminal and type following command, give same topic name which have given in producer:

.\bin\windows\kafka-console-consumer.bat --topic {topic name} --from-beginning --bootstrap-server localhost:9092

```
C:\Windows\System32\cmd.exe - \bin\windows\kafka-console-consumer.bat --topic preExams --from-beginning --bootstrap-server localhost:9092 — X

Microsoft Windows [Version 10.0.19945.3570]
(c) Microsoft Corporation. All rights reserved.

D:\kafka>.\bin\windows\kafka-console-consumer.bat --topic preExams --from-beginning --bootstrap-server localhost:9092
```

Now set producer and consumer side by side, type something in producer and hit enter after some millisecond you can see same thing in consumer terminal.