

HPC_17

April 16, 2024

1 Assignment 17

Implement following CUDA programs

1. to print hello message on the screen using kernal function
2. to add two vectors of size 100 and 20000 abd analyze the performance comparison between cpu and gpu processing
3. to multiply two matrix of size 20 X 20 and 1024 X 1024 analyze the performance comparison between cpu and gpu processing
4. to obtain CUDA device information and print the output

```
[1]: !pip install pycuda
```

Collecting pycuda

Downloading pycuda-2024.1.tar.gz (1.7 MB)

1.7/1.7 MB

29.7 MB/s eta 0:00:00

Installing build dependencies ... done

Getting requirements to build wheel ... done

Preparing metadata (pyproject.toml) ... done

Collecting pytools>=2011.2 (from pycuda)

Downloading pytools-2024.1.1-py2.py3-none-any.whl (85 kB)

85.1/85.1 kB

13.5 MB/s eta 0:00:00

Requirement already satisfied: appdirs>=1.4.0 in

/usr/local/lib/python3.10/dist-packages (from pycuda) (1.4.4)

Collecting mako (from pycuda)

Downloading Mako-1.3.3-py3-none-any.whl (78 kB)

78.8/78.8 kB

11.6 MB/s eta 0:00:00

Requirement already satisfied: platformdirs>=2.2.0 in

/usr/local/lib/python3.10/dist-packages (from pytools>=2011.2->pycuda) (4.2.0)

Requirement already satisfied: typing-extensions>=4.0 in

/usr/local/lib/python3.10/dist-packages (from pytools>=2011.2->pycuda) (4.11.0)

Requirement already satisfied: MarkupSafe>=0.9.2 in

/usr/local/lib/python3.10/dist-packages (from mako->pycuda) (2.1.5)

Building wheels for collected packages: pycuda

Building wheel for pycuda (pyproject.toml) ... done

Created wheel for pycuda: filename=pycuda-2024.1-cp310-cp310-linux_x86_64.whl
size=661204

sha256=c6d745905ea408f8a73a79e4656b609ac0f9dced62157e4766697031d2a03d17

Stored in directory: /root/.cache/pip/wheels/12/34/d2/9a349255a4eca3a486d82c79d21e138ce2ccd90f414d9d72b8
Successfully built pycuda
Installing collected packages: pytools, mako, pycuda
Successfully installed mako-1.3.3 pycuda-2024.1 pytools-2024.1.1

1.1 1. To print hello message on the screen using kernal function

```
[10]: %%writefile hello_1_1.cu

#include <stdio.h>

__global__ void cuda_hello_1_1() {
    printf("Hello World from GPU with grid dimension (1, 1) and block dimension (1, 1)!\n");
}

int main() {
    cuda_hello_1_1<<<1,1>>>();
    cudaDeviceSynchronize(); // Make sure all GPU work is done before exiting
    return 0;
}
```

Writing hello_1_1.cu

```
[11]: !nvcc -o hello_1_1 hello_1_1.cu
```

```
[12]: !./hello_1_1
```

Hello World from GPU with grid dimension (1, 1) and block dimension (1, 1)!

1.2 2. To add two vectors of size 100 and 20000 and analyze the performance comparison between cpu and gpu processing

1.2.1 GPU

```
[2]: import numpy as np
import pycuda.driver as cuda
import pycuda.autotinit
from pycuda.compiler import SourceModule
import time
```

```
[3]: cuda_kernel_code = """
__global__ void vector_add(float *a, float *b, float *c, int n) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n) {
        c[i] = a[i] + b[i];
    }
}
```

```

"""

cuda_module = SourceModule(cuda_kernel_code)

vector_add_cuda = cuda_module.get_function("vector_add")

def vector_add_gpu(a, b):
    n = a.size

    a_gpu = cuda.mem_alloc(a.nbytes)
    b_gpu = cuda.mem_alloc(b.nbytes)
    c_gpu = cuda.mem_alloc(b.nbytes)

    cuda.memcpy_htod(a_gpu, a)
    cuda.memcpy_htod(b_gpu, b)

    block_dim = (256, 1, 1)
    grid_dim = ((n + block_dim[0] - 1) // block_dim[0], 1)

    start_time = time.time()
    vector_add_cuda(a_gpu, b_gpu, c_gpu, np.int32(n), block=block_dim,
↪grid=grid_dim)

    cuda.Context.synchronize()

    end_time = time.time()

    c = np.empty_like(a)
    cuda.memcpy_dtoh(c, c_gpu)

    return c, end_time - start_time

vector_size_1 = 100
vector_size_2 = 20000
a = np.random.randn(vector_size_2).astype(np.float32)
b = np.random.randn(vector_size_2).astype(np.float32)

result_gpu1, gpu_time1 = vector_add_gpu(a[:vector_size_1], b[:vector_size_1])
result_gpu2, gpu_time2 = vector_add_gpu(a[:vector_size_2], b[:vector_size_2])

print("Vector addition of size", vector_size_1, "on GPU took", gpu_time1,
↪"seconds.")
print("Vector addition of size", vector_size_2, "on GPU took", gpu_time2,
↪"seconds.")

```

Vector addition of size 100 on GPU took 0.0006387233734130859 seconds.

Vector addition of size 20000 on GPU took 5.435943603515625e-05 seconds.

1.2.2 CPU

```
[4]: import numpy as np
import time
```

```
[5]: def vector_add_cpu(a, b):
    start_time = time.time()
    result = a + b
    end_time = time.time()
    return result, end_time - start_time
```

```
[6]: vector_size_1 = 100
vector_size_2 = 20000
a = np.random.randn(vector_size_2).astype(np.float32)
b = np.random.randn(vector_size_2).astype(np.float32)

result_cpu1, cpu_time1 = vector_add_cpu(a[:vector_size_1], b[:vector_size_1])
result_cpu2, cpu_time2 = vector_add_cpu(a[:vector_size_2], b[:vector_size_2])
print("Vector addition of size", vector_size_1, "on CPU took", cpu_time1, "seconds.")
print("Vector addition of size", vector_size_2, "on CPU took", cpu_time2, "seconds.")
```

Vector addition of size 100 on CPU took 2.384185791015625e-05 seconds.

Vector addition of size 20000 on CPU took 2.0742416381835938e-05 seconds.

1.3 3. To multiply two matrix of size 20 X 20 and 1024 X 1024 analyze the performance comparison between cpu and gpu processing

1.3.1 GPU

```
[7]: def matrix_multiply_gpu(a, b):
    cuda_code = """
    __global__ void matrix_multiply(float *a, float *b, float *c, int n) {
        int row = blockIdx.y * blockDim.y + threadIdx.y;
        int col = blockIdx.x * blockDim.x + threadIdx.x;

        if (row < n && col < n) {
            float sum = 0.0;
            for (int i = 0; i < n; ++i) {
                sum += a[row * n + i] * b[i * n + col];
            }
            c[row * n + col] = sum;
        }
    }
    """
```

```

mod = SourceModule(cuda_code)

matrix_multiply_cuda = mod.get_function("matrix_multiply")

a_gpu = cuda.mem_alloc(a.nbytes)
b_gpu = cuda.mem_alloc(b.nbytes)
c_gpu = cuda.mem_alloc(a.nbytes)

cuda.memcpy_htod(a_gpu, a)
cuda.memcpy_htod(b_gpu, b)

block_size = (16, 16, 1)
grid_size = ((a.shape[1] + block_size[0] - 1) // block_size[0], (a.shape[0]
↪ + block_size[1] - 1) // block_size[1], 1)

matrix_multiply_cuda(a_gpu, b_gpu, c_gpu, np.int32(a.shape[0]),
↪ block=block_size, grid=grid_size)

c = np.empty_like(a)
cuda.memcpy_dtoh(c, c_gpu)

return c

def generate_random_matrix(rows, cols):
    return np.random.rand(rows, cols).astype(np.float32)

def measure_time(matrix_size, func, *args):
    start_time = time.time()
    result = func(*args)
    end_time = time.time()
    return result, end_time - start_time

matrix_sizes = [(20, 20), (1024, 1024)]

for size in matrix_sizes:
    print(f"\nMatrix size: {size}")
    a = generate_random_matrix(*size)
    b = generate_random_matrix(*size)

    gpu_result, gpu_time = measure_time(size, matrix_multiply_gpu, a, b)
    print(f"GPU time: {gpu_time:.6f} seconds")

```

Matrix size: (20, 20)
GPU time: 0.437754 seconds

Matrix size: (1024, 1024)

GPU time: 0.013317 seconds

1.3.2 CPU

```
[8]: def matrix_multiply_cpu(a, b):
    result = np.zeros((a.shape[0], b.shape[1]), dtype=np.float32)
    for i in range(a.shape[0]):
        for j in range(b.shape[1]):
            for k in range(a.shape[1]):
                result[i, j] += a[i, k] * b[k, j]
    return result

def generate_random_matrix(rows, cols):
    return np.random.rand(rows, cols).astype(np.float32)

def measure_time(matrix_size, func, *args):
    start_time = time.time()
    result = func(*args)
    end_time = time.time()
    return result, end_time - start_time

matrix_sizes = [(20, 20), (1024, 1024)]

for size in matrix_sizes:
    print(f"\nMatrix size: {size}")
    a = generate_random_matrix(*size)
    b = generate_random_matrix(*size)

    # CPU matrix multiplication
    cpu_result, cpu_time = measure_time(size, matrix_multiply_cpu, a, b)
    print(f"CPU time: {cpu_time:.6f} seconds")
```

Matrix size: (20, 20)

CPU time: 0.009491 seconds

Matrix size: (1024, 1024)

CPU time: 718.713489 seconds

2 4. To obtain CUDA device information and print the output

```
[ ]: import pycuda.driver as cuda

cuda.init()

num_devices = cuda.Device.count()
```

```

print("Number of CUDA devices:", num_devices)

for i in range(num_devices):
    device = cuda.Device(i)
    print("\nCUDA Device:", i)
    print("  Name:", device.name())
    print("  Compute Capability:", device.compute_capability())
    print("  Total Memory:", device.total_memory() / (1024 ** 3), "GB")
    print("  Max Threads per Block:", device.max_threads_per_block)
    print("  Multiprocessor Count:", device.multiprocessor_count)
    print("  Clock Rate:", device.clock_rate / 1e6, "GHz")

```

Number of CUDA devices: 1

CUDA Device: 0

Name: Tesla T4

Compute Capability: (7, 5)

Total Memory: 14.74810791015625 GB

Max Threads per Block: 1024

Multiprocessor Count: 40

Clock Rate: 1.59 GHz

[9]: !nvidia-smi

Tue Apr 16 08:54:22 2024

```

+-----+
+-----+
| NVIDIA-SMI 535.104.05           Driver Version: 535.104.05   CUDA Version:
12.2           |
+-----+-----+-----+
+-----+
| GPU  Name                Persistence-M | Bus-Id        Disp.A | Volatile
Uncorr. ECC |
| Fan  Temp   Perf          Pwr:Usage/Cap |      Memory-Usage | GPU-Util
Compute M.  |
|                                           |                   |
MIG M.       |
+=====+
=====|
|   0   Tesla T4                          Off | 00000000:00:04.0 Off |
0 |
| N/A    58C    P0              28W / 70W |      103MiB / 15360MiB |      0%
Default |
|                                           |                   |
N/A |
+-----+-----+-----+
+-----+

```

```

+-----+
-----+
| Processes:
|
| GPU  GI  CI          PID  Type  Process name                      GPU
Memory |
|          ID  ID
Usage   |
|=====
=====|
+-----+
-----+

```