

# HPC\_assignment\_17

April 16, 2024

## 1 HPC assignment 17

```
[ ]: !nvidia-smi
```

Tue Apr 16 07:06:05 2024

```
+-----+
+-----+
| NVIDIA-SMI 535.104.05                  Driver Version: 535.104.05   CUDA Version:
12.2          |
+-----+-----+-----+
+-----+
| GPU  Name                       Persistence-M | Bus-Id        Disp.A | Volatile
Uncorr. ECC |
| Fan   Temp   Perf           Pwr:Usage/Cap |      Memory-Usage | GPU-Util
Compute M. |
|              |              |              |
MIG M. |
|=====+=====+=====+
=====|
|   0   Tesla T4                       Off | 00000000:00:04.0 Off |
0 |
| N/A    40C    P8              9W / 70W |      0MiB / 15360MiB |      0%
Default |
|              |              |              |
N/A |
+-----+-----+-----+
+-----+
+-----+
| Processes:
|
| GPU   GI    CI          PID    Type    Process name                        GPU
Memory |
|      ID    ID
Usage   |
|=====+=====+=====+
=====|
```

```
| No running processes found
```

```
|
```

```
+-----+  
-----+
```

## 1.1 1. To print hello message on the screen using kernal function

```
[ ]: %%writefile hello_1_1.cu  
  
#include <stdio.h>  
  
__global__ void cuda_hello_1_1() {  
    printf("Hello World from GPU with grid dimension (1, 1) and block dimension_  
    ↪(1, 1)!\n");  
}  
  
int main() {  
    cuda_hello_1_1<<<1,1>>>();  
    cudaDeviceSynchronize(); // Make sure all GPU work is done before exiting  
    return 0;  
}
```

Writing hello\_1\_1.cu

```
[ ]: !nvcc -o hello_1_1 hello_1_1.cu
```

```
[ ]: !./hello_1_1
```

Hello World from GPU with grid dimension (1, 1) and block dimension (1, 1)!

## 1.2 2. To add two vectors of size 100 and 20000 and analyze the performance comparison between cpu and gpu processing

### 1.2.1 GPU

```
[ ]: !pip install pycuda
```

Collecting pycuda

Downloading pycuda-2024.1.tar.gz (1.7 MB)

1.7/1.7 MB

12.0 MB/s eta 0:00:00

Installing build dependencies ... done

Getting requirements to build wheel ... done

Preparing metadata (pyproject.toml) ... done

Collecting pytools>=2011.2 (from pycuda)

Downloading pytools-2024.1.1-py2.py3-none-any.whl (85 kB)

85.1/85.1 kB

12.0 MB/s eta 0:00:00

Requirement already satisfied: appdirs>=1.4.0 in

```

/usr/local/lib/python3.10/dist-packages (from pycuda) (1.4.4)
Collecting mako (from pycuda)
  Downloading Mako-1.3.3-py3-none-any.whl (78 kB)
    78.8/78.8 kB
10.9 MB/s eta 0:00:00
Requirement already satisfied: platformdirs>=2.2.0 in
/usr/local/lib/python3.10/dist-packages (from pytools>=2011.2->pycuda) (4.2.0)
Requirement already satisfied: typing-extensions>=4.0 in
/usr/local/lib/python3.10/dist-packages (from pytools>=2011.2->pycuda) (4.11.0)
Requirement already satisfied: MarkupSafe>=0.9.2 in
/usr/local/lib/python3.10/dist-packages (from mako->pycuda) (2.1.5)
Building wheels for collected packages: pycuda
  Building wheel for pycuda (pyproject.toml) ... done
  Created wheel for pycuda: filename=pycuda-2024.1-cp310-cp310-linux_x86_64.whl
  size=661204
  sha256=51efb7c5582dd86e48b9404a05e0a366352406f4840bf4dc162fe9a89aa2ad1c
  Stored in directory: /root/.cache/pip/wheels/12/34/d2/9a349255a4eca3a486d82c79
d21e138ce2ccd90f414d9d72b8
Successfully built pycuda
Installing collected packages: pytools, mako, pycuda
Successfully installed mako-1.3.3 pycuda-2024.1 pytools-2024.1.1

```

```

[ ]: import numpy as np
import pycuda.driver as cuda
import pycuda.autoinit
from pycuda.compiler import SourceModule
import time

```

```

[ ]: # CUDA kernel function to add two vectors
cuda_kernel_code = """
__global__ void vector_add(float *a, float *b, float *c, int n) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n) {
        c[i] = a[i] + b[i];
    }
}
"""

```

```

[ ]: # Compile the CUDA kernel code
cuda_module = SourceModule(cuda_kernel_code)

# Get a reference to the CUDA kernel function
vector_add_cuda = cuda_module.get_function("vector_add")

```

```

[ ]: def vector_add_gpu(a, b):
    n = a.size

    # Create device arrays

```

```

a_gpu = cuda.mem_alloc(a.nbytes)
b_gpu = cuda.mem_alloc(b.nbytes)
c_gpu = cuda.mem_alloc(b.nbytes)

# Copy data to device
cuda.memcpy_htod(a_gpu, a)
cuda.memcpy_htod(b_gpu, b)

# Define block and grid dimensions
block_dim = (256, 1, 1)
grid_dim = ((n + block_dim[0] - 1) // block_dim[0], 1)

start_time = time.time()

# Launch the CUDA kernel
vector_add_cuda(a_gpu, b_gpu, c_gpu, np.int32(n), block=block_dim,
↪grid=grid_dim)

# Synchronize threads to ensure all output is calculated
cuda.Context.synchronize()

end_time = time.time()

# Copy result back to host
c = np.empty_like(a)
cuda.memcpy_dtoh(c, c_gpu)

return c, end_time - start_time

```

```

[ ]: vector_size_1 = 100
vector_size_2 = 20000
a = np.random.randn(vector_size_2).astype(np.float32)
b = np.random.randn(vector_size_2).astype(np.float32)

# Perform vector addition on GPU
result_gpu1, gpu_time1 = vector_add_gpu(a[:vector_size_1], b[:vector_size_1])
result_gpu2, gpu_time2 = vector_add_gpu(a[:vector_size_2], b[:vector_size_2])

[ ]: print("Vector addition of size", vector_size_1, "on GPU took", gpu_time1,
↪"seconds.")

print("Vector addition of size", vector_size_2, "on GPU took", gpu_time2,
↪"seconds.")

```

Vector addition of size 100 on GPU took 0.0007643699645996094 seconds.  
Vector addition of size 20000 on GPU took 6.818771362304688e-05 seconds.

### 1.2.2 CPU

```
[ ]: import numpy as np
import time

[ ]: def vector_add_cpu(a, b):
    start_time = time.time()
    result = a + b
    end_time = time.time()
    return result, end_time - start_time

[ ]: vector_size_1 = 100
vector_size_2 = 20000
a = np.random.randn(vector_size_2).astype(np.float32)
b = np.random.randn(vector_size_2).astype(np.float32)

# Perform vector addition on CPU
result_cpu1, cpu_time1 = vector_add_cpu(a[:vector_size_1], b[:vector_size_1])
result_cpu2, cpu_time2 = vector_add_cpu(a[:vector_size_2], b[:vector_size_2])

[ ]: print("Vector addition of size", vector_size_1, "on CPU took", cpu_time1,↵
↵"seconds.")
print("Vector addition of size", vector_size_2, "on CPU took", cpu_time2,↵
↵"seconds.")
```

Vector addition of size 100 on CPU took 2.3365020751953125e-05 seconds.

Vector addition of size 20000 on CPU took 1.9311904907226562e-05 seconds.

- Vector addition of size 100 on CPU took 2.384185791015625e-05 seconds.
- Vector addition of size 20000 on CPU took 1.9788742065429688e-05 seconds.
- Vector addition of size 100 on GPU took 0.0007691383361816406 seconds.
- Vector addition of size 20000 on GPU took 7.128715515136719e-05 seconds.

## 1.3 3. To multiply two matrix of size 20 X 20 and 1024 X 1024 analyze the performance comparison between cpu and gpu processing

### 1.3.1 GPU

```
[ ]: def matrix_multiply_gpu(a, b):
    # Define CUDA kernel code for matrix multiplication
    cuda_code = """
    __global__ void matrix_multiply(float *a, float *b, float *c, int n) {
        int row = blockIdx.y * blockDim.y + threadIdx.y;
        int col = blockIdx.x * blockDim.x + threadIdx.x;

        if (row < n && col < n) {
            float sum = 0.0;
            for (int i = 0; i < n; ++i) {
                sum += a[row * n + i] * b[i * n + col];
            }
            c[row * n + col] = sum;
        }
    }
    """
```

```

        }
        c[row * n + col] = sum;
    }
}
"""

# Compile CUDA kernel code
mod = SourceModule(cuda_code)

# Get kernel function
matrix_multiply_cuda = mod.get_function("matrix_multiply")

# Allocate memory on device
a_gpu = cuda.mem_alloc(a.nbytes)
b_gpu = cuda.mem_alloc(b.nbytes)
c_gpu = cuda.mem_alloc(a.nbytes)

# Copy input matrices to device
cuda.memcpy_htod(a_gpu, a)
cuda.memcpy_htod(b_gpu, b)

# Define grid and block dimensions
block_size = (16, 16, 1)
grid_size = ((a.shape[1] + block_size[0] - 1) // block_size[0], (a.shape[0]
↪ + block_size[1] - 1) // block_size[1], 1)

# Call CUDA kernel
matrix_multiply_cuda(a_gpu, b_gpu, c_gpu, np.int32(a.shape[0]),
↪ block_size, grid=grid_size)

# Copy result back to host
c = np.empty_like(a)
cuda.memcpy_dtoh(c, c_gpu)

return c

```

```

[ ]: # Function to generate random matrices
def generate_random_matrix(rows, cols):
    return np.random.rand(rows, cols).astype(np.float32)

```

```

[ ]: # Function to measure time taken for matrix multiplication
def measure_time(matrix_size, func, *args):
    start_time = time.time()
    result = func(*args)
    end_time = time.time()
    return result, end_time - start_time

```

```
[ ]: # Sizes of matrices to be multiplied
matrix_sizes = [(20, 20), (1024, 1024)]
```

```
[ ]: for size in matrix_sizes:
    print(f"\nMatrix size: {size}")
    a = generate_random_matrix(*size)
    b = generate_random_matrix(*size)

    # GPU matrix multiplication
    gpu_result, gpu_time = measure_time(size, matrix_multiply_gpu, a, b)
    print(f"GPU time: {gpu_time:.6f} seconds")
```

Matrix size: (20, 20)  
GPU time: 0.428407 seconds

Matrix size: (1024, 1024)  
GPU time: 0.018636 seconds

### 1.3.2 CPU

```
[ ]: # CPU matrix multiplication
def matrix_multiply_cpu(a, b):
    result = np.zeros((a.shape[0], b.shape[1]), dtype=np.float32)
    for i in range(a.shape[0]):
        for j in range(b.shape[1]):
            for k in range(a.shape[1]):
                result[i, j] += a[i, k] * b[k, j]
    return result
```

```
[ ]: # Function to generate random matrices
def generate_random_matrix(rows, cols):
    return np.random.rand(rows, cols).astype(np.float32)
```

```
[ ]: # Function to measure time taken for matrix multiplication
def measure_time(matrix_size, func, *args):
    start_time = time.time()
    result = func(*args)
    end_time = time.time()
    return result, end_time - start_time
```

```
[ ]: # Sizes of matrices to be multiplied
matrix_sizes = [(20, 20), (1024, 1024)]
```

```
[ ]: for size in matrix_sizes:
    print(f"\nMatrix size: {size}")
    a = generate_random_matrix(*size)
    b = generate_random_matrix(*size)
```

```
# CPU matrix multiplication
cpu_result, cpu_time = measure_time(size, matrix_multiply_cpu, a, b)
print(f"CPU time: {cpu_time:.6f} seconds")
```

Matrix size: (20, 20)  
CPU time: 0.004824 seconds

Matrix size: (1024, 1024)  
CPU time: 704.230331 seconds

- CPU Time for 1024: 0.12308359146118164 seconds
- CPU Time for 20: 0.0019140243530273438 seconds
- Matrix size: (20, 20)
- GPU time: 0.703994 seconds
- Matrix size: (1024, 1024)
- GPU time: 0.014648 seconds

#### 1.4 4. To obtain CUDA device information and print the output

```
[ ]: import pycuda.driver as cuda

# Initialize PyCUDA
cuda.init()

# Get the number of CUDA devices
num_devices = cuda.Device.count()

print("Number of CUDA devices:", num_devices)

# Iterate over each CUDA device and print its properties
for i in range(num_devices):
    device = cuda.Device(i)
    print("\nCUDA Device:", i)
    print("  Name:", device.name())
    print("  Compute Capability:", device.compute_capability())
    print("  Total Memory:", device.total_memory() / (1024 ** 3), "GB")
    print("  Max Threads per Block:", device.max_threads_per_block)
    print("  Multiprocessor Count:", device.multiprocessor_count)
    print("  Clock Rate:", device.clock_rate / 1e6, "GHz")
```

Number of CUDA devices: 1

CUDA Device: 0  
Name: Tesla T4  
Compute Capability: (7, 5)



Total Memory: 14.74810791015625 GB  
Max Threads per Block: 1024  
Multiprocessor Count: 40  
Clock Rate: 1.59 GHz